

Bogumiła Walkowiak bogumila.walkowiak@grenoble-inp.org

Joachim Małowski joachim-kajetan.makowski@grenoble-inp.org

Intelligent Systems: Reasoning and Recognition

Recognizing Digits using Neural Networks

1. Introduction

The MNIST (Modified National Institute of Standards and Technology) dataset is a large collection of handwritten digits composed of 60,000 training images and 10,000 test images. The black and white images from NIST were normalized to fit into a 28x28 pixel bounding box and anti-aliased, which introduced gray-scale levels. Our task was to design and evaluate neural network architectures that can recognize hand-drawn digits using the grayscale this data set.

2. Data preparation

First of all, we downloaded MNIST data. We decided to combine train and test set provided by MNIST dataset and then we split data into training set 90% and a test set 10%. In the further part of the project, we'll also create a validation set so the final split of the data will look like this: training data 80%, validating data 10% and testing data 10%.

```
In [86]: import numpy as np
import tensorflow.compat.v1.keras.backend as K
import tensorflow as tf
from tensorflow import keras

from tensorflow.keras import layers
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.metrics import f1_score, roc_curve, auc

from sklearn.metrics import confusion_matrix, plot_confusion_matrix
#physical_devices = tf.config.list_physical_devices('GPU')
#tf.config.experimental.set_memory_growth(physical_devices[0], True)
```

```
In [95]: (x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()
```


```
In [96]: X = np.concatenate((x_train, x_test))
y = np.concatenate([y_train, y_test])
```

```
In [97]: train_ratio = 0.9
test_ratio = 0.1
```

```
x_train, x_test, y_train, y_test = train_test_split(X, y, test_size = te
```

```
In [98]: plt.imshow(x_train[0], cmap='gray')
```

```
Out[98]: <matplotlib.image.AxesImage at 0x7f3cfc7c7b90>
```



```
In [99]: x_train = x_train.astype("float32") / 255
x_test = x_test.astype("float32") / 255

# images have shape (28, 28, 1)
x_train = np.expand_dims(x_train, -1)
x_test = np.expand_dims(x_test, -1)
```

```
In [100]: y_train = keras.utils.to_categorical(y_train, 10)
y_test = keras.utils.to_categorical(y_test, 10)
```

3. Creating neural networks

We decided to create a function. Thanks to it we will be able to write less code. Function trains model provided by argument of function, prints model's loss, accuracy, precision, recall and AUC for each digit and plots a history of training.

```
In [60]: def predict_model(model, callbacks = [], batch_size=128, epochs = 4, lr=0.
adam = keras.optimizers.Adam(lr=lr)
model.compile(loss="categorical_crossentropy", optimizer=adam, metri

history = model.fit(x_train, y_train, batch_size=batch_size, epochs=
score = model.evaluate(x_test, y_test, verbose=0)
y_pred = model.predict(x_test)
print("Test loss:", score[0])
print("Test accuracy:", score[1])
print("Test precision:", score[2])
print("Test recall:", score[3])

y_pred = np.argmax(y_pred, axis=1)
y_test1 = np.argmax(y_test, axis=1)

print("Test f1 score:", f1_score(y_test1, y_pred, average='micro'))
for i in range(10):
    temp_pred = [1 if x==i else 0 for x in y_pred]
```

```

temp_test = [1 if x==i else 0 for x in y_test1]
fpr, tpr, thresholds =roc_curve(temp_test,temp_pred)

print("Test AUC for digit:",i, auc(fpr, tpr))

# summarize history for accuracy
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()
# summarize history for loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()

```

We added an instance of EarlyStopping class, which provides us a mechanism of stopping algorithm before the whole training process is done. When 3 epochs are not achieving a better result (in our example higher validation accuracy) then our training is stopped and we restore the best model.

```

In [61]: # simple early stopping
es = keras.callbacks.EarlyStopping(monitor='val_accuracy', mode='max', v

```

Basic Fully Connected Multi-layer Network

The first network we have created is basic fully connected mutli-layer network:

```

In [48]: model_fc = keras.Sequential([
    layers.Dense(32, activation="relu",input_shape=(28,28,1)),
    layers.Dense(64, activation="relu"),

    layers.Flatten(),
    layers.Dense(128, activation="relu"),
    layers.Dropout(.25),
    layers.Dense(10, activation="softmax")
])
model_fc.summary()
predict_model(model_fc, [es], epochs=100)

```

Model: "sequential_20"

Layer (type)	Output Shape	Param #
dense_59 (Dense)	(None, 28, 28, 32)	64
dense_60 (Dense)	(None, 28, 28, 64)	2112
flatten_20 (Flatten)	(None, 50176)	0
dense_61 (Dense)	(None, 128)	6422656
dropout_19 (Dropout)	(None, 128)	0

dense_62 (Dense) (None, 10) 1290

=====
Total params: 6,426,122
Trainable params: 6,426,122
Non-trainable params: 0

Epoch 1/100

439/439 [=====] - 6s 13ms/step - loss: 0.5066 -
accuracy: 0.8487 - precision: 0.9149 - recall: 0.7786 - val_loss: 0.1463
- val_accuracy: 0.9586 - val_precision: 0.9667 - val_recall: 0.9504

Epoch 2/100

439/439 [=====] - 5s 12ms/step - loss: 0.1643 -
accuracy: 0.9510 - precision: 0.9601 - recall: 0.9434 - val_loss: 0.1087
- val_accuracy: 0.9683 - val_precision: 0.9707 - val_recall: 0.9661

Epoch 3/100

439/439 [=====] - 5s 12ms/step - loss: 0.1254 -
accuracy: 0.9611 - precision: 0.9680 - recall: 0.9554 - val_loss: 0.0970
- val_accuracy: 0.9729 - val_precision: 0.9770 - val_recall: 0.9698

Epoch 4/100

439/439 [=====] - 5s 12ms/step - loss: 0.1044 -
accuracy: 0.9666 - precision: 0.9716 - recall: 0.9629 - val_loss: 0.0932
- val_accuracy: 0.9736 - val_precision: 0.9756 - val_recall: 0.9711

Epoch 5/100

439/439 [=====] - 5s 12ms/step - loss: 0.0893 -
accuracy: 0.9704 - precision: 0.9740 - recall: 0.9670 - val_loss: 0.0962
- val_accuracy: 0.9720 - val_precision: 0.9748 - val_recall: 0.9710

Epoch 6/100

439/439 [=====] - 5s 12ms/step - loss: 0.0824 -
accuracy: 0.9726 - precision: 0.9757 - recall: 0.9695 - val_loss: 0.0853
- val_accuracy: 0.9773 - val_precision: 0.9801 - val_recall: 0.9746

Epoch 7/100

439/439 [=====] - 5s 12ms/step - loss: 0.0750 -
accuracy: 0.9763 - precision: 0.9789 - recall: 0.9733 - val_loss: 0.0796
- val_accuracy: 0.9769 - val_precision: 0.9796 - val_recall: 0.9755

Epoch 8/100

439/439 [=====] - 5s 12ms/step - loss: 0.0664 -
accuracy: 0.9776 - precision: 0.9805 - recall: 0.9752 - val_loss: 0.0869
- val_accuracy: 0.9766 - val_precision: 0.9787 - val_recall: 0.9755


Epoch 9/100

439/439 [=====] - 5s 12ms/step - loss: 0.0621 -
accuracy: 0.9788 - precision: 0.9806 - recall: 0.9770 - val_loss: 0.0832
- val_accuracy: 0.9766 - val_precision: 0.9783 - val_recall: 0.9758

Restoring model weights from the end of the best epoch.

Epoch 00009: early stopping

Test loss: 0.09832347929477692
Test accuracy: 0.973714292049408
Test precision: 0.9774295687675476
Test recall: 0.9712857007980347
Test f1 score: 0.9737142857142858
Test AUC for digit: 0 0.9907399351644153
Test AUC for digit: 1 0.9923331284991935
Test AUC for digit: 2 0.9870421148228989
Test AUC for digit: 3 0.9851853814963031
Test AUC for digit: 4 0.9845406523282492
Test AUC for digit: 5 0.9781305833322657
Test AUC for digit: 6 0.9906349206349208
Test AUC for digit: 7 0.9875151552516612
Test AUC for digit: 8 0.974920634920635
Test AUC for digit: 9 0.9815057646170433



This is basic model achieves about 97,5% accuracy on test set. It is made of 2 hidden layers with reasonable number of units. Training this model is quite fast (on my laptop it was 5s per epoch, using GPU). As we see in plots our model started to overfits, because validation accuracy and loss was staying on the same level, while train accuracy was growing and loss was decreasing.

Next, we wanted to demonstrate the effect of changing various parameters of the network.

Different number of layers

```
In [13]: model_fc_small = keras.Sequential([
            layers.Dense(32, activation="relu", input_shape=(28,28,1)),

            layers.Flatten(),
            layers.Dense(64, activation="relu"),
            layers.Dropout(.25),
            layers.Dense(10, activation="softmax")
        ])
model_fc_small.summary()
predict_model(model_fc_small, [es], epochs=100)
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		

dense_4 (Dense)	(None, 28, 28, 32)	64
flatten_1 (Flatten)	(None, 25088)	0
dense_5 (Dense)	(None, 64)	1605696
dropout_1 (Dropout)	(None, 64)	0
dense_6 (Dense)	(None, 10)	650

=====
Total params: 1,606,410
Trainable params: 1,606,410
Non-trainable params: 0

Epoch 1/100
439/439 [=====] - 11s 22ms/step - loss: 0.5918 - accuracy: 0.8171 - precision: 0.9070 - recall: 0.7388 - val_loss: 0.1764 - val_accuracy: 0.9483 - val_precision: 0.9603 - val_recall: 0.9381

Epoch 2/100
439/439 [=====] - 9s 21ms/step - loss: 0.2418 - accuracy: 0.9282 - precision: 0.9433 - recall: 0.9133 - val_loss: 0.1424 - val_accuracy: 0.9557 - val_precision: 0.9640 - val_recall: 0.9494

Epoch 3/100
439/439 [=====] - 10s 22ms/step - loss: 0.1958 - accuracy: 0.9404 - precision: 0.9512 - recall: 0.9299 - val_loss: 0.1180 - val_accuracy: 0.9649 - val_precision: 0.9711 - val_recall: 0.9587

Epoch 4/100
439/439 [=====] - 9s 21ms/step - loss: 0.1628 - accuracy: 0.9478 - precision: 0.9577 - recall: 0.9405 - val_loss: 0.1205 - val_accuracy: 0.9635 - val_precision: 0.9701 - val_recall: 0.9589

Epoch 5/100
439/439 [=====] - 9s 22ms/step - loss: 0.1513 - accuracy: 0.9516 - precision: 0.9591 - recall: 0.9447 - val_loss: 0.1079 - val_accuracy: 0.9680 - val_precision: 0.9741 - val_recall: 0.9644

Epoch 6/100
439/439 [=====] - 9s 22ms/step - loss: 0.1431 - accuracy: 0.9555 - precision: 0.9634 - recall: 0.9489 - val_loss: 0.1074 - val_accuracy: 0.9690 - val_precision: 0.9747 - val_recall: 0.9674

Epoch 7/100
439/439 [=====] - 10s 22ms/step - loss: 0.1351 - accuracy: 0.9565 - precision: 0.9637 - recall: 0.9509 - val_loss: 0.1093 - val_accuracy: 0.9694 - val_precision: 0.9741 - val_recall: 0.9646

Epoch 8/100
439/439 [=====] - 10s 22ms/step - loss: 0.1319 - accuracy: 0.9587 - precision: 0.9651 - recall: 0.9528 - val_loss: 0.1047 - val_accuracy: 0.9703 - val_precision: 0.9741 - val_recall: 0.9674

Epoch 9/100
439/439 [=====] - 10s 22ms/step - loss: 0.1188 - accuracy: 0.9614 - precision: 0.9674 - recall: 0.9562 - val_loss: 0.1044 - val_accuracy: 0.9706 - val_precision: 0.9747 - val_recall: 0.9670

Epoch 10/100
439/439 [=====] - 10s 22ms/step - loss: 0.1194 - accuracy: 0.9610 - precision: 0.9666 - recall: 0.9566 - val_loss: 0.1080 - val_accuracy: 0.9724 - val_precision: 0.9749 - val_recall: 0.9697



Epoch 11/100
439/439 [=====] - 10s 22ms/step - loss: 0.1116 - accuracy: 0.9636 - precision: 0.9691 - recall: 0.9591 - val_loss: 0.1011 - val_accuracy: 0.9732 - val_precision: 0.9760 - val_recall: 0.9693

Epoch 12/100
439/439 [=====] - 10s 22ms/step - loss: 0.1030 - accuracy: 0.9653 - precision: 0.9705 - recall: 0.9611 - val_loss: 0.1011 - val_accuracy: 0.9722 - val_precision: 0.9743 - val_recall: 0.9698

Epoch 13/100
439/439 [=====] - 10s 22ms/step - loss: 0.1039 - accuracy: 0.9663 - precision: 0.9707 - recall: 0.9623 - val_loss: 0.1028 - val_accuracy: 0.9729 - val_precision: 0.9764 - val_recall: 0.9713

Epoch 14/100
439/439 [=====] - 10s 22ms/step - loss: 0.1029 -

accuracy: 0.9668 - precision: 0.9710 - recall: 0.9625 - val_loss: 0.0977
- val_accuracy: 0.9730 - val_precision: 0.9772 - val_recall: 0.9717
Epoch 00014: early stopping
Test loss: 0.1229480430483818
Test accuracy: 0.9679999947547913
Test precision: 0.9723661541938782
Test recall: 0.9651428461074829
Test f1 score: 0.968
Test AUC for digit: 0 0.9900292592326402
Test AUC for digit: 1 0.9906050941195285
Test AUC for digit: 2 0.9857744868717027
Test AUC for digit: 3 0.9805160602125461
Test AUC for digit: 4 0.981489857731353
Test AUC for digit: 5 0.9728804648288869
Test AUC for digit: 6 0.9881746031746033
Test AUC for digit: 7 0.9815794883823452
Test AUC for digit: 8 0.9792063492063492
Test AUC for digit: 9 0.9701437036094706



```
In [14]: model_fc_large = keras.Sequential([
    layers.Dense(32, activation="relu", input_shape=(28,28,1)),
    layers.Dense(64, activation="relu"),
    layers.Flatten(),
    layers.Dense(4096, activation="relu"),
    layers.Dense(1024, activation="relu"),
    layers.Dense(64, activation="relu"),
    layers.Dropout(.25),
    layers.Dense(10, activation="softmax")
])
```

```
model_fc_large.summary()  
predict_model(model_fc_large, [es], epochs=100)
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_7 (Dense)	(None, 28, 28, 32)	64
dense_8 (Dense)	(None, 28, 28, 64)	2112
flatten_2 (Flatten)	(None, 50176)	0
dense_9 (Dense)	(None, 4096)	205524992
dense_10 (Dense)	(None, 1024)	4195328
dense_11 (Dense)	(None, 64)	65600
dropout_2 (Dropout)	(None, 64)	0
dense_12 (Dense)	(None, 10)	650
Total params: 209,788,746		
Trainable params: 209,788,746		
Non-trainable params: 0		

Epoch 1/100

439/439 [=====] - 55s 120ms/step - loss: 0.5493
- accuracy: 0.8354 - precision: 0.9044 - recall: 0.7825 - val_loss: 0.123
6 - val_accuracy: 0.9652 - val_precision: 0.9722 - val_recall: 0.9605

Epoch 2/100

439/439 [=====] - 33s 75ms/step - loss: 0.1123 -
accuracy: 0.9679 - precision: 0.9741 - recall: 0.9620 - val_loss: 0.0866
- val_accuracy: 0.9762 - val_precision: 0.9798 - val_recall: 0.9717

Epoch 3/100

439/439 [=====] - 33s 75ms/step - loss: 0.0735 -
accuracy: 0.9787 - precision: 0.9823 - recall: 0.9749 - val_loss: 0.0845
- val_accuracy: 0.9747 - val_precision: 0.9812 - val_recall: 0.9709

Epoch 4/100

439/439 [=====] - 33s 75ms/step - loss: 0.0509 -
accuracy: 0.9847 - precision: 0.9870 - recall: 0.9826 - val_loss: 0.0801
- val_accuracy: 0.9778 - val_precision: 0.9823 - val_recall: 0.9752

Epoch 5/100

439/439 [=====] - 33s 75ms/step - loss: 0.0464 -
accuracy: 0.9860 - precision: 0.9879 - recall: 0.9842 - val_loss: 0.0887
- val_accuracy: 0.9781 - val_precision: 0.9806 - val_recall: 0.9762

Epoch 6/100

439/439 [=====] - 33s 75ms/step - loss: 0.0356 -
accuracy: 0.9898 - precision: 0.9909 - recall: 0.9882 - val_loss: 0.0897
- val_accuracy: 0.9804 - val_precision: 0.9828 - val_recall: 0.9785

Epoch 7/100

439/439 [=====] - 33s 75ms/step - loss: 0.0276 -
accuracy: 0.9921 - precision: 0.9929 - recall: 0.9913 - val_loss: 0.0870
- val_accuracy: 0.9821 - val_precision: 0.9832 - val_recall: 0.9811

Epoch 8/100

439/439 [=====] - 33s 75ms/step - loss: 0.0234 -
accuracy: 0.9928 - precision: 0.9935 - recall: 0.9920 - val_loss: 0.0820
- val_accuracy: 0.9804 - val_precision: 0.9823 - val_recall: 0.9776

Epoch 9/100

439/439 [=====] - 33s 75ms/step - loss: 0.0215 -
accuracy: 0.9939 - precision: 0.9947 - recall: 0.9931 - val_loss: 0.1025
- val_accuracy: 0.9801 - val_precision: 0.9812 - val_recall: 0.9786

Epoch 10/100


439/439 [=====] - 32s 73ms/step - loss: 0.0220 -
accuracy: 0.9939 - precision: 0.9943 - recall: 0.9933 - val_loss: 0.0775
- val_accuracy: 0.9828 - val_precision: 0.9849 - val_recall: 0.9821


Epoch 11/100


```

439/439 [=====] - 34s 77ms/step - loss: 0.0177 -
accuracy: 0.9949 - precision: 0.9953 - recall: 0.9944 - val_loss: 0.0929
- val_accuracy: 0.9843 - val_precision: 0.9850 - val_recall: 0.9831
Epoch 12/100
439/439 [=====] - 34s 79ms/step - loss: 0.0138 -
accuracy: 0.9958 - precision: 0.9962 - recall: 0.9955 - val_loss: 0.0812
- val_accuracy: 0.9815 - val_precision: 0.9828 - val_recall: 0.9808
Epoch 13/100
439/439 [=====] - 34s 77ms/step - loss: 0.0136 -
accuracy: 0.9963 - precision: 0.9965 - recall: 0.9960 - val_loss: 0.0883
- val_accuracy: 0.9824 - val_precision: 0.9834 - val_recall: 0.9818
Epoch 14/100
439/439 [=====] - 34s 77ms/step - loss: 0.0117 -
accuracy: 0.9965 - precision: 0.9967 - recall: 0.9963 - val_loss: 0.0809
- val_accuracy: 0.9844 - val_precision: 0.9858 - val_recall: 0.9844
Epoch 15/100
439/439 [=====] - 34s 78ms/step - loss: 0.0107 -
accuracy: 0.9971 - precision: 0.9972 - recall: 0.9968 - val_loss: 0.1161
- val_accuracy: 0.9805 - val_precision: 0.9811 - val_recall: 0.9798
Epoch 16/100
439/439 [=====] - 33s 74ms/step - loss: 0.0152 -
accuracy: 0.9959 - precision: 0.9960 - recall: 0.9955 - val_loss: 0.1062
- val_accuracy: 0.9815 - val_precision: 0.9827 - val_recall: 0.9814
Epoch 17/100
439/439 [=====] - 34s 78ms/step - loss: 0.0097 -
accuracy: 0.9971 - precision: 0.9973 - recall: 0.9970 - val_loss: 0.0868
- val_accuracy: 0.9830 - val_precision: 0.9841 - val_recall: 0.9824
Epoch 00017: early stopping
Test loss: 0.09394106268882751
Test accuracy: 0.9814285635948181
Test precision: 0.9819768071174622
Test recall: 0.9807142615318298
Test f1 score: 0.9814285714285714
Test AUC for digit: 0 0.9988155401137082
Test AUC for digit: 1 0.9947057004231706
Test AUC for digit: 2 0.9863037767355775
Test AUC for digit: 3 0.9898547027800599
Test AUC for digit: 4 0.9875124205303294
Test AUC for digit: 5 0.9887902209336588
Test AUC for digit: 6 0.9899999999999999
Test AUC for digit: 7 0.9909183219141172
Test AUC for digit: 8 0.9815079365079366
Test AUC for digit: 9 0.987926026320382

```





Firstly, we tried different numbers of hidden layers. With 1 hidden layer the model the model was achieving around 96,5% on test set. The model is underfitted because this number of layers is not enough to explain the complexity of our data.

Model with 4 hidden layers achieved 98,1% of accuracy but the training time was pretty long (34s per epoch). That is because this model had to find weights for over 200,000,000 parameters (compering to 1,600,000 of params for model with 1 hidden layer). We can assume, that after second epoch our model is overfitted because the difference between validation and train loss and accuracy are high.

Different number of units per layer

```
In [15]: model_fc = keras.Sequential([
    layers.Dense(10, activation="relu", input_shape=(28,28,1)),
    layers.Dense(20, activation="relu"),

    layers.Flatten(),
    layers.Dense(40, activation="relu"),
    layers.Dropout(.25),
    layers.Dense(10, activation="softmax")
])
model_fc.summary()
predict_model(model_fc, [es], epochs=100)
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
dense_13 (Dense)	(None, 28, 28, 10)	20
dense_14 (Dense)	(None, 28, 28, 20)	220
flatten_3 (Flatten)	(None, 15680)	0
dense_15 (Dense)	(None, 40)	627240
dropout_3 (Dropout)	(None, 40)	0
dense_16 (Dense)	(None, 10)	410
Total params: 627,890		
Trainable params: 627,890		
Non-trainable params: 0		

Epoch 1/100
439/439 [=====] - 4s 8ms/step - loss: 0.6942 - accuracy: 0.7885 - precision: 0.8959 - recall: 0.6762 - val_loss: 0.2011 - val_accuracy: 0.9401 - val_precision: 0.9534 - val_recall: 0.9297

Epoch 2/100
439/439 [=====] - 3s 7ms/step - loss: 0.2938 - accuracy: 0.9101 - precision: 0.9330 - recall: 0.8907 - val_loss: 0.1577 - val_accuracy: 0.9567 - val_precision: 0.9660 - val_recall: 0.9476

Epoch 3/100
439/439 [=====] - 3s 7ms/step - loss: 0.2441 - accuracy: 0.9252 - precision: 0.9418 - recall: 0.9100 - val_loss: 0.1455 - val_accuracy: 0.9571 - val_precision: 0.9666 - val_recall: 0.9519

Epoch 4/100
439/439 [=====] - 3s 7ms/step - loss: 0.2146 - accuracy: 0.9344 - precision: 0.9488 - recall: 0.9216 - val_loss: 0.1342 - val_accuracy: 0.9613 - val_precision: 0.9710 - val_recall: 0.9534

Epoch 5/100
439/439 [=====] - 3s 7ms/step - loss: 0.2027 - accuracy: 0.9370 - precision: 0.9503 - recall: 0.9245 - val_loss: 0.1228 - val_accuracy: 0.9639 - val_precision: 0.9711 - val_recall: 0.9589

Epoch 6/100
439/439 [=====] - 3s 7ms/step - loss: 0.1886 - accuracy: 0.9395 - precision: 0.9515 - recall: 0.9284 - val_loss: 0.1218 - val_accuracy: 0.9613 - val_precision: 0.9689 - val_recall: 0.9582

Epoch 7/100
439/439 [=====] - 3s 8ms/step - loss: 0.1840 - accuracy: 0.9408 - precision: 0.9528 - recall: 0.9316 - val_loss: 0.1211 - val_accuracy: 0.9649 - val_precision: 0.9713 - val_recall: 0.9612

Epoch 8/100
439/439 [=====] - 3s 7ms/step - loss: 0.1801 - accuracy: 0.9424 - precision: 0.9536 - recall: 0.9319 - val_loss: 0.1177 - val_accuracy: 0.9646 - val_precision: 0.9717 - val_recall: 0.9613

Epoch 9/100
439/439 [=====] - 3s 7ms/step - loss: 0.1690 - accuracy: 0.9457 - precision: 0.9563 - recall: 0.9362 - val_loss: 0.1225 - val_accuracy: 0.9657 - val_precision: 0.9714 - val_recall: 0.9615

Epoch 10/100
439/439 [=====] - 3s 7ms/step - loss: 0.1685 - accuracy: 0.9465 - precision: 0.9569 - recall: 0.9375 - val_loss: 0.1151 - val_accuracy: 0.9675 - val_precision: 0.9729 - val_recall: 0.9644

Epoch 11/100
439/439 [=====] - 3s 7ms/step - loss: 0.1615 - accuracy: 0.9484 - precision: 0.9589 - recall: 0.9397 - val_loss: 0.1208 - val_accuracy: 0.9658 - val_precision: 0.9719 - val_recall: 0.9596

Epoch 12/100
439/439 [=====] - 3s 7ms/step - loss: 0.1573 - accuracy: 0.9494 - precision: 0.9584 - recall: 0.9422 - val_loss: 0.1155 - val_accuracy: 0.9685 - val_precision: 0.9716 - val_recall: 0.9641

Epoch 13/100
439/439 [=====] - 3s 7ms/step - loss: 0.1518 - accuracy: 0.9498 - precision: 0.9592 - recall: 0.9414 - val_loss: 0.1165 - val_accuracy: 0.9658 - val_precision: 0.9719 - val_recall: 0.9616

Epoch 14/100
439/439 [=====] - 3s 7ms/step - loss: 0.1459 - accuracy: 0.9534 - precision: 0.9620 - recall: 0.9461 - val_loss: 0.1136 - val_accuracy: 0.9670 - val_precision: 0.9723 - val_recall: 0.9641


Epoch 15/100
439/439 [=====] - 3s 7ms/step - loss: 0.1485 - accuracy: 0.9509 - precision: 0.9603 - recall: 0.9437 - val_loss: 0.1178 - val_accuracy: 0.9667 - val_precision: 0.9704 - val_recall: 0.9641

Epoch 00015: early stopping
Test loss: 0.12980997562408447
Test accuracy: 0.9620000123977661
Test precision: 0.9678164124488831
Test recall: 0.9580000042915344
Test f1 score: 0.962
Test AUC for digit: 0 0.9925116601919346

```

Test AUC for digit: 1 0.9898799893821776
Test AUC for digit: 2 0.9862624960357879
Test AUC for digit: 3 0.9697247767206424
Test AUC for digit: 4 0.9822645042790503
Test AUC for digit: 5 0.9736597292505758
Test AUC for digit: 6 0.9884920634920635
Test AUC for digit: 7 0.9744171471575807
Test AUC for digit: 8 0.9637301587301589
Test AUC for digit: 9 0.9673563983253335

```



In this situation we trained a model with small number of units in each layer. The model didn't achieve it's best. We can see, that train accuracy is much lower than validation accuracy. It is caused by insufficient number of units, so that our model decided to choose higher accuracy in validation data at the expense of accuracy on whole data.

```

In [16]: model_fc = keras.Sequential([
    layers.Dense(100, activation="relu", input_shape=(28,28,1)),
    layers.Dense(200, activation="relu"),

    layers.Flatten(),
    layers.Dense(400, activation="relu"),
    layers.Dropout(.25),
    layers.Dense(10, activation="softmax")
])
model_fc.summary()
predict_model(model_fc, [es], epochs=100)

```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
dense_17 (Dense)	(None, 28, 28, 100)	200
dense_18 (Dense)	(None, 28, 28, 200)	20200
flatten_4 (Flatten)	(None, 156800)	0
dense_19 (Dense)	(None, 400)	62720400
dropout_4 (Dropout)	(None, 400)	0
dense_20 (Dense)	(None, 10)	4010
Total params: 62,744,810		
Trainable params: 62,744,810		
Non-trainable params: 0		

Epoch 1/100

439/439 [=====] - 20s 45ms/step - loss: 0.3810 - accuracy: 0.8801 - precision: 0.9206 - recall: 0.8528 - val_loss: 0.1177 - val_accuracy: 0.9678 - val_precision: 0.9722 - val_recall: 0.9632

Epoch 2/100

439/439 [=====] - 19s 44ms/step - loss: 0.1207 - accuracy: 0.9633 - precision: 0.9690 - recall: 0.9591 - val_loss: 0.0924 - val_accuracy: 0.9750 - val_precision: 0.9782 - val_recall: 0.9726

Epoch 3/100

439/439 [=====] - 19s 44ms/step - loss: 0.0939 - accuracy: 0.9698 - precision: 0.9735 - recall: 0.9668 - val_loss: 0.0772 - val_accuracy: 0.9772 - val_precision: 0.9801 - val_recall: 0.9756

Epoch 4/100

439/439 [=====] - 19s 44ms/step - loss: 0.0664 - accuracy: 0.9785 - precision: 0.9809 - recall: 0.9767 - val_loss: 0.0717 - val_accuracy: 0.9804 - val_precision: 0.9828 - val_recall: 0.9789

Epoch 5/100

439/439 [=====] - 19s 44ms/step - loss: 0.0588 - accuracy: 0.9814 - precision: 0.9831 - recall: 0.9800 - val_loss: 0.0762 - val_accuracy: 0.9789 - val_precision: 0.9824 - val_recall: 0.9771

Epoch 6/100

439/439 [=====] - 19s 44ms/step - loss: 0.0463 - accuracy: 0.9847 - precision: 0.9860 - recall: 0.9832 - val_loss: 0.0695 - val_accuracy: 0.9808 - val_precision: 0.9831 - val_recall: 0.9798

Epoch 7/100

439/439 [=====] - 19s 44ms/step - loss: 0.0436 - accuracy: 0.9851 - precision: 0.9862 - recall: 0.9839 - val_loss: 0.0762 - val_accuracy: 0.9797 - val_precision: 0.9813 - val_recall: 0.9784

Epoch 8/100

439/439 [=====] - 19s 44ms/step - loss: 0.0367 - accuracy: 0.9878 - precision: 0.9884 - recall: 0.9864 - val_loss: 0.0631 - val_accuracy: 0.9823 - val_precision: 0.9836 - val_recall: 0.9808

Epoch 9/100

439/439 [=====] - 19s 44ms/step - loss: 0.0343 - accuracy: 0.9876 - precision: 0.9886 - recall: 0.9868 - val_loss: 0.0663 - val_accuracy: 0.9807 - val_precision: 0.9830 - val_recall: 0.9782

Epoch 10/100

439/439 [=====] - 19s 44ms/step - loss: 0.0300 - accuracy: 0.9894 - precision: 0.9902 - recall: 0.9889 - val_loss: 0.0815 - val_accuracy: 0.9828 - val_precision: 0.9831 - val_recall: 0.9820

Epoch 11/100

439/439 [=====] - 19s 44ms/step - loss: 0.0253 - accuracy: 0.9917 - precision: 0.9922 - recall: 0.9912 - val_loss: 0.0702 - val_accuracy: 0.9834 - val_precision: 0.9847 - val_recall: 0.9830



Epoch 12/100

439/439 [=====] - 19s 44ms/step - loss: 0.0224 - accuracy: 0.9931 - precision: 0.9933 - recall: 0.9928 - val_loss: 0.0760 - val_accuracy: 0.9818 - val_precision: 0.9831 - val_recall: 0.9811

```

Epoch 13/100
439/439 [=====] - 19s 44ms/step - loss: 0.0239 -
accuracy: 0.9915 - precision: 0.9921 - recall: 0.9910 - val_loss: 0.0783
- val_accuracy: 0.9798 - val_precision: 0.9818 - val_recall: 0.9789
Epoch 14/100
439/439 [=====] - 19s 44ms/step - loss: 0.0202 -
accuracy: 0.9927 - precision: 0.9931 - recall: 0.9925 - val_loss: 0.0734
- val_accuracy: 0.9808 - val_precision: 0.9828 - val_recall: 0.9795
Epoch 00014: early stopping
Test loss: 0.08533725142478943
Test accuracy: 0.9795714020729065
Test precision: 0.9812401533126831
Test recall: 0.978857159614563
Test f1 score: 0.9795714285714285
Test AUC for digit: 0 0.993617156085807
Test AUC for digit: 1 0.9955623103018043
Test AUC for digit: 2 0.9869502393935624
Test AUC for digit: 3 0.9921123419818196
Test AUC for digit: 4 0.9906739460192185
Test AUC for digit: 5 0.9834615107924303
Test AUC for digit: 6 0.993095238095238
Test AUC for digit: 7 0.990522564434301
Test AUC for digit: 8 0.976031746031746
Test AUC for digit: 9 0.9826922857371192

```



In this model we see that it's overfitting after third epoch. It is caused by too high number of units.

Different learning rate

In [49]:

```
model_fc_01 = keras.Sequential([
    layers.Dense(32, activation="relu", input_shape=(28,28,1)),
    layers.Dense(64, activation="relu"),

    layers.Flatten(),
    layers.Dense(128, activation="relu"),
    layers.Dropout(.25),
    layers.Dense(10, activation="softmax")
])
model_fc_01.summary()
predict_model(model_fc_01,[es], epochs=100, lr=0.05)
```

Model: "sequential_21"

Layer (type)	Output Shape	Param #
dense_63 (Dense)	(None, 28, 28, 32)	64
dense_64 (Dense)	(None, 28, 28, 64)	2112
flatten_21 (Flatten)	(None, 50176)	0
dense_65 (Dense)	(None, 128)	6422656
dropout_20 (Dropout)	(None, 128)	0
dense_66 (Dense)	(None, 10)	1290

Total params: 6,426,122
Trainable params: 6,426,122
Non-trainable params: 0

Epoch 1/100

439/439 [=====] - 6s 13ms/step - loss: 2.5641 - accuracy: 0.5516 - precision: 0.7483 - recall: 0.4283 - val_loss: 0.4579 - val_accuracy: 0.8913 - val_precision: 0.9146 - val_recall: 0.8734

Epoch 2/100

439/439 [=====] - 5s 12ms/step - loss: 0.6291 - accuracy: 0.8000 - precision: 0.8749 - recall: 0.7328 - val_loss: 0.2885 - val_accuracy: 0.9248 - val_precision: 0.9536 - val_recall: 0.8866

Epoch 3/100

439/439 [=====] - 5s 12ms/step - loss: 0.5281 - accuracy: 0.8320 - precision: 0.8874 - recall: 0.7802 - val_loss: 0.2769 - val_accuracy: 0.9264 - val_precision: 0.9468 - val_recall: 0.9092

Epoch 4/100

439/439 [=====] - 5s 12ms/step - loss: 0.5307 - accuracy: 0.8348 - precision: 0.8777 - recall: 0.7971 - val_loss: 0.2822 - val_accuracy: 0.9258 - val_precision: 0.9504 - val_recall: 0.8993

Epoch 5/100

439/439 [=====] - 5s 12ms/step - loss: 0.6051 - accuracy: 0.8107 - precision: 0.8614 - recall: 0.7684 - val_loss: 0.2608 - val_accuracy: 0.9299 - val_precision: 0.9458 - val_recall: 0.9134

Epoch 6/100

439/439 [=====] - 5s 12ms/step - loss: 0.4968 - accuracy: 0.8454 - precision: 0.8837 - recall: 0.8143 - val_loss: 0.2928 - val_accuracy: 0.9188 - val_precision: 0.9445 - val_recall: 0.8965

Epoch 7/100

439/439 [=====] - 5s 12ms/step - loss: 0.5996 - accuracy: 0.8170 - precision: 0.8669 - recall: 0.7753 - val_loss: 0.3559 - val_accuracy: 0.9036 - val_precision: 0.9382 - val_recall: 0.8711



Epoch 8/100

439/439 [=====] - 5s 12ms/step - loss: 0.5959 - accuracy: 0.8130 - precision: 0.8761 - recall: 0.7515 - val_loss: 0.3750 - val_accuracy: 0.9023 - val_precision: 0.9415 - val_recall: 0.8524

Restoring model weights from the end of the best epoch.

Epoch 00008: early stopping

Test loss: 0.28543806076049805
 Test accuracy: 0.9227142930030823
 Test precision: 0.942704439163208
 Test recall: 0.9072856903076172
 Test f1 score: 0.9227142857142857
 Test AUC for digit: 0 0.9789563324393538
 Test AUC for digit: 1 0.979840545957394
 Test AUC for digit: 2 0.9702686971098221
 Test AUC for digit: 3 0.927925323986961
 Test AUC for digit: 4 0.9539853844616494
 Test AUC for digit: 5 0.9262426101687231
 Test AUC for digit: 6 0.9711111111111111
 Test AUC for digit: 7 0.960725329011793
 Test AUC for digit: 8 0.9488095238095239
 Test AUC for digit: 9 0.9511993241314747



We took our first model and decided to train it with different learning rates. With learning rate 0.05 we received very bad results (accuracy around 92%). The scores are so bad because our optimizer did not find good weights, because it had to change values with too big "jump".

```

In [39]: model_fc_00001 = keras.Sequential([
    layers.Dense(32, activation="relu", input_shape=(28,28,1)),
    layers.Dense(64, activation="relu"),

    layers.Flatten(),
    layers.Dense(128, activation="relu"),
    layers.Dropout(.25),
  
```



```

        layers.Dense(10, activation="softmax")
    ])
    model_fc_00001.summary()
    predict_model(model_fc_00001,[es], epochs=100, lr = 0.00001)

```

Model: "sequential_15"

Layer (type)	Output Shape	Param #
dense_46 (Dense)	(None, 28, 28, 32)	64
dense_47 (Dense)	(None, 28, 28, 64)	2112
flatten_15 (Flatten)	(None, 50176)	0
dense_48 (Dense)	(None, 128)	6422656
dropout_14 (Dropout)	(None, 128)	0
dense_49 (Dense)	(None, 10)	1290
Total params: 6,426,122		
Trainable params: 6,426,122		
Non-trainable params: 0		

Epoch 1/100

439/439 [=====] - 6s 13ms/step - loss: 1.8177 - accuracy: 0.5851 - precision: 0.7578 - recall: 0.0475 - val_loss: 0.8030 - val_accuracy: 0.8551 - val_precision: 0.9845 - val_recall: 0.5235

Epoch 2/100

439/439 [=====] - 5s 12ms/step - loss: 0.7526 - accuracy: 0.8255 - precision: 0.9614 - recall: 0.5864 - val_loss: 0.5059 - val_accuracy: 0.8835 - val_precision: 0.9580 - val_recall: 0.7645

Epoch 3/100

439/439 [=====] - 5s 12ms/step - loss: 0.5323 - accuracy: 0.8581 - precision: 0.9378 - recall: 0.7558 - val_loss: 0.4048 - val_accuracy: 0.8981 - val_precision: 0.9491 - val_recall: 0.8320

Epoch 4/100

439/439 [=====] - 5s 12ms/step - loss: 0.4437 - accuracy: 0.8764 - precision: 0.9324 - recall: 0.8142 - val_loss: 0.3561 - val_accuracy: 0.9055 - val_precision: 0.9460 - val_recall: 0.8603

Epoch 5/100

439/439 [=====] - 5s 12ms/step - loss: 0.3928 - accuracy: 0.8891 - precision: 0.9336 - recall: 0.8411 - val_loss: 0.3259 - val_accuracy: 0.9117 - val_precision: 0.9436 - val_recall: 0.8758

Epoch 6/100

439/439 [=====] - 5s 12ms/step - loss: 0.3637 - accuracy: 0.8953 - precision: 0.9334 - recall: 0.8587 - val_loss: 0.3050 - val_accuracy: 0.9153 - val_precision: 0.9434 - val_recall: 0.8876

Epoch 7/100

439/439 [=====] - 5s 12ms/step - loss: 0.3479 - accuracy: 0.8986 - precision: 0.9315 - recall: 0.8666 - val_loss: 0.2869 - val_accuracy: 0.9190 - val_precision: 0.9459 - val_recall: 0.8964

Epoch 8/100

439/439 [=====] - 5s 12ms/step - loss: 0.3248 - accuracy: 0.9051 - precision: 0.9349 - recall: 0.8769 - val_loss: 0.2723 - val_accuracy: 0.9227 - val_precision: 0.9477 - val_recall: 0.9000

Epoch 9/100

439/439 [=====] - 5s 12ms/step - loss: 0.3000 - accuracy: 0.9150 - precision: 0.9403 - recall: 0.8894 - val_loss: 0.2609 - val_accuracy: 0.9258 - val_precision: 0.9487 - val_recall: 0.9040

Epoch 10/100

439/439 [=====] - 5s 12ms/step - loss: 0.2921 - accuracy: 0.9155 - precision: 0.9409 - recall: 0.8923 - val_loss: 0.2493 - val_accuracy: 0.9293 - val_precision: 0.9492 - val_recall: 0.9094

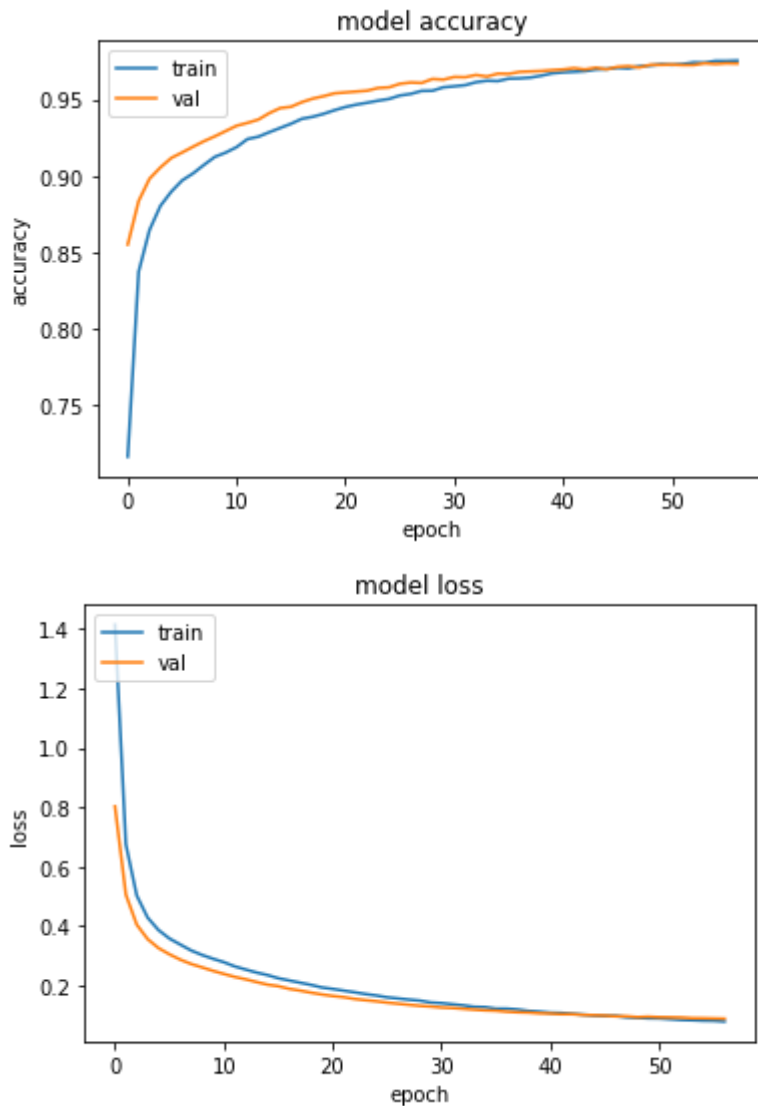
Epoch 11/100

439/439 [=====] - 5s 12ms/step - loss: 0.2765 -

accuracy: 0.9198 - precision: 0.9424 - recall: 0.8979 - val_loss: 0.2389
- val_accuracy: 0.9328 - val_precision: 0.9516 - val_recall: 0.9134
Epoch 12/100
439/439 [=====] - 5s 12ms/step - loss: 0.2636 -
accuracy: 0.9249 - precision: 0.9456 - recall: 0.9054 - val_loss: 0.2288
- val_accuracy: 0.9346 - val_precision: 0.9539 - val_recall: 0.9189
Epoch 13/100
439/439 [=====] - 5s 12ms/step - loss: 0.2534 -
accuracy: 0.9258 - precision: 0.9459 - recall: 0.9079 - val_loss: 0.2207
- val_accuracy: 0.9368 - val_precision: 0.9539 - val_recall: 0.9216
Epoch 14/100
439/439 [=====] - 5s 12ms/step - loss: 0.2447 -
accuracy: 0.9276 - precision: 0.9473 - recall: 0.9091 - val_loss: 0.2116
- val_accuracy: 0.9410 - val_precision: 0.9564 - val_recall: 0.9264
Epoch 15/100
439/439 [=====] - 5s 12ms/step - loss: 0.2372 -
accuracy: 0.9315 - precision: 0.9506 - recall: 0.9155 - val_loss: 0.2028
- val_accuracy: 0.9443 - val_precision: 0.9583 - val_recall: 0.9296
Epoch 16/100
439/439 [=====] - 5s 12ms/step - loss: 0.2272 -
accuracy: 0.9334 - precision: 0.9515 - recall: 0.9177 - val_loss: 0.1975
- val_accuracy: 0.9452 - val_precision: 0.9590 - val_recall: 0.9326
Epoch 17/100
439/439 [=====] - 5s 12ms/step - loss: 0.2178 -
accuracy: 0.9368 - precision: 0.9532 - recall: 0.9223 - val_loss: 0.1886
- val_accuracy: 0.9482 - val_precision: 0.9601 - val_recall: 0.9369
Epoch 18/100
439/439 [=====] - 5s 12ms/step - loss: 0.2111 -
accuracy: 0.9389 - precision: 0.9525 - recall: 0.9257 - val_loss: 0.1827
- val_accuracy: 0.9505 - val_precision: 0.9626 - val_recall: 0.9390
Epoch 19/100
439/439 [=====] - 5s 12ms/step - loss: 0.2045 -
accuracy: 0.9398 - precision: 0.9540 - recall: 0.9263 - val_loss: 0.1758
- val_accuracy: 0.9521 - val_precision: 0.9638 - val_recall: 0.9418
Epoch 20/100
439/439 [=====] - 5s 12ms/step - loss: 0.1990 -
accuracy: 0.9418 - precision: 0.9564 - recall: 0.9299 - val_loss: 0.1696
- val_accuracy: 0.9540 - val_precision: 0.9645 - val_recall: 0.9443
Epoch 21/100
439/439 [=====] - 5s 12ms/step - loss: 0.1892 -
accuracy: 0.9436 - precision: 0.9570 - recall: 0.9323 - val_loss: 0.1645
- val_accuracy: 0.9545 - val_precision: 0.9649 - val_recall: 0.9455
Epoch 22/100
439/439 [=====] - 5s 12ms/step - loss: 0.1896 -
accuracy: 0.9450 - precision: 0.9588 - recall: 0.9329 - val_loss: 0.1605
- val_accuracy: 0.9551 - val_precision: 0.9655 - val_recall: 0.9479
Epoch 23/100
439/439 [=====] - 5s 12ms/step - loss: 0.1779 -
accuracy: 0.9476 - precision: 0.9589 - recall: 0.9375 - val_loss: 0.1546
- val_accuracy: 0.9557 - val_precision: 0.9679 - val_recall: 0.9489
Epoch 24/100
439/439 [=====] - 5s 12ms/step - loss: 0.1749 -
accuracy: 0.9487 - precision: 0.9613 - recall: 0.9386 - val_loss: 0.1504
- val_accuracy: 0.9574 - val_precision: 0.9670 - val_recall: 0.9509
Epoch 25/100
439/439 [=====] - 5s 12ms/step - loss: 0.1666 -
accuracy: 0.9498 - precision: 0.9622 - recall: 0.9410 - val_loss: 0.1468
- val_accuracy: 0.9579 - val_precision: 0.9681 - val_recall: 0.9509
Epoch 26/100
439/439 [=====] - 5s 12ms/step - loss: 0.1642 -
accuracy: 0.9515 - precision: 0.9622 - recall: 0.9420 - val_loss: 0.1422
- val_accuracy: 0.9602 - val_precision: 0.9699 - val_recall: 0.9534
Epoch 27/100
439/439 [=====] - 5s 12ms/step - loss: 0.1584 -
accuracy: 0.9539 - precision: 0.9639 - recall: 0.9446 - val_loss: 0.1384
- val_accuracy: 0.9610 - val_precision: 0.9695 - val_recall: 0.9548
Epoch 28/100
439/439 [=====] - 5s 12ms/step - loss: 0.1521 -

accuracy: 0.9556 - precision: 0.9661 - recall: 0.9471 - val_loss: 0.1349
- val_accuracy: 0.9608 - val_precision: 0.9707 - val_recall: 0.9545
Epoch 29/100
439/439 [=====] - 5s 12ms/step - loss: 0.1494 -
accuracy: 0.9548 - precision: 0.9648 - recall: 0.9473 - val_loss: 0.1314
- val_accuracy: 0.9633 - val_precision: 0.9719 - val_recall: 0.9567
Epoch 30/100
439/439 [=====] - 5s 12ms/step - loss: 0.1435 -
accuracy: 0.9581 - precision: 0.9669 - recall: 0.9507 - val_loss: 0.1290
- val_accuracy: 0.9629 - val_precision: 0.9714 - val_recall: 0.9569
Epoch 31/100
439/439 [=====] - 5s 12ms/step - loss: 0.1403 -
accuracy: 0.9578 - precision: 0.9669 - recall: 0.9508 - val_loss: 0.1257
- val_accuracy: 0.9646 - val_precision: 0.9718 - val_recall: 0.9586
Epoch 32/100
439/439 [=====] - 5s 12ms/step - loss: 0.1395 -
accuracy: 0.9596 - precision: 0.9680 - recall: 0.9526 - val_loss: 0.1241
- val_accuracy: 0.9645 - val_precision: 0.9722 - val_recall: 0.9595
Epoch 33/100
439/439 [=====] - 5s 12ms/step - loss: 0.1311 -
accuracy: 0.9620 - precision: 0.9695 - recall: 0.9547 - val_loss: 0.1213
- val_accuracy: 0.9659 - val_precision: 0.9721 - val_recall: 0.9609
Epoch 34/100
439/439 [=====] - 5s 12ms/step - loss: 0.1297 -
accuracy: 0.9622 - precision: 0.9695 - recall: 0.9555 - val_loss: 0.1186
- val_accuracy: 0.9649 - val_precision: 0.9721 - val_recall: 0.9605
Epoch 35/100
439/439 [=====] - 5s 12ms/step - loss: 0.1279 -
accuracy: 0.9621 - precision: 0.9696 - recall: 0.9551 - val_loss: 0.1162
- val_accuracy: 0.9668 - val_precision: 0.9726 - val_recall: 0.9629
Epoch 36/100
439/439 [=====] - 5s 12ms/step - loss: 0.1146 -
accuracy: 0.9653 - precision: 0.9719 - recall: 0.9593 - val_loss: 0.1150
- val_accuracy: 0.9667 - val_precision: 0.9724 - val_recall: 0.9616
Epoch 37/100
439/439 [=====] - 5s 12ms/step - loss: 0.1249 -
accuracy: 0.9629 - precision: 0.9696 - recall: 0.9567 - val_loss: 0.1117
- val_accuracy: 0.9680 - val_precision: 0.9730 - val_recall: 0.9629
Epoch 38/100
439/439 [=====] - 5s 12ms/step - loss: 0.1214 -
accuracy: 0.9626 - precision: 0.9702 - recall: 0.9565 - val_loss: 0.1097
- val_accuracy: 0.9683 - val_precision: 0.9736 - val_recall: 0.9635
Epoch 39/100
439/439 [=====] - 5s 12ms/step - loss: 0.1171 -
accuracy: 0.9645 - precision: 0.9716 - recall: 0.9592 - val_loss: 0.1087
- val_accuracy: 0.9685 - val_precision: 0.9730 - val_recall: 0.9638
Epoch 40/100
439/439 [=====] - 5s 12ms/step - loss: 0.1124 -
accuracy: 0.9678 - precision: 0.9735 - recall: 0.9615 - val_loss: 0.1062
- val_accuracy: 0.9691 - val_precision: 0.9735 - val_recall: 0.9649
Epoch 41/100
439/439 [=====] - 5s 12ms/step - loss: 0.1069 -
accuracy: 0.9684 - precision: 0.9746 - recall: 0.9629 - val_loss: 0.1053
- val_accuracy: 0.9696 - val_precision: 0.9739 - val_recall: 0.9651
Epoch 42/100
439/439 [=====] - 5s 12ms/step - loss: 0.1087 -
accuracy: 0.9674 - precision: 0.9734 - recall: 0.9625 - val_loss: 0.1033
- val_accuracy: 0.9704 - val_precision: 0.9747 - val_recall: 0.9657
Epoch 43/100
439/439 [=====] - 5s 12ms/step - loss: 0.1096 -
accuracy: 0.9669 - precision: 0.9732 - recall: 0.9618 - val_loss: 0.1032
- val_accuracy: 0.9696 - val_precision: 0.9744 - val_recall: 0.9655
Epoch 44/100
439/439 [=====] - 5s 12ms/step - loss: 0.1057 -
accuracy: 0.9685 - precision: 0.9743 - recall: 0.9640 - val_loss: 0.1008
- val_accuracy: 0.9707 - val_precision: 0.9743 - val_recall: 0.9671
Epoch 45/100
439/439 [=====] - 5s 12ms/step - loss: 0.0991 -

accuracy: 0.9693 - precision: 0.9747 - recall: 0.9651 - val_loss: 0.0997
- val_accuracy: 0.9697 - val_precision: 0.9745 - val_recall: 0.9668
Epoch 46/100
439/439 [=====] - 5s 12ms/step - loss: 0.0953 -
accuracy: 0.9718 - precision: 0.9766 - recall: 0.9676 - val_loss: 0.0985
- val_accuracy: 0.9713 - val_precision: 0.9746 - val_recall: 0.9681
Epoch 47/100
439/439 [=====] - 5s 12ms/step - loss: 0.0971 -
accuracy: 0.9703 - precision: 0.9753 - recall: 0.9661 - val_loss: 0.0962
- val_accuracy: 0.9717 - val_precision: 0.9759 - val_recall: 0.9696
Epoch 48/100
439/439 [=====] - 5s 12ms/step - loss: 0.0926 -
accuracy: 0.9709 - precision: 0.9765 - recall: 0.9662 - val_loss: 0.0961
- val_accuracy: 0.9709 - val_precision: 0.9753 - val_recall: 0.9684
Epoch 49/100
439/439 [=====] - 5s 12ms/step - loss: 0.0916 -
accuracy: 0.9721 - precision: 0.9774 - recall: 0.9678 - val_loss: 0.0933
- val_accuracy: 0.9726 - val_precision: 0.9769 - val_recall: 0.9703
Epoch 50/100
439/439 [=====] - 5s 12ms/step - loss: 0.0899 -
accuracy: 0.9725 - precision: 0.9768 - recall: 0.9677 - val_loss: 0.0946
- val_accuracy: 0.9726 - val_precision: 0.9762 - val_recall: 0.9688
Epoch 51/100
439/439 [=====] - 5s 12ms/step - loss: 0.0869 -
accuracy: 0.9727 - precision: 0.9777 - recall: 0.9684 - val_loss: 0.0929
- val_accuracy: 0.9727 - val_precision: 0.9762 - val_recall: 0.9701
Epoch 52/100
439/439 [=====] - 5s 12ms/step - loss: 0.0876 -
accuracy: 0.9723 - precision: 0.9773 - recall: 0.9680 - val_loss: 0.0924
- val_accuracy: 0.9724 - val_precision: 0.9763 - val_recall: 0.9700
Epoch 53/100
439/439 [=====] - 5s 12ms/step - loss: 0.0830 -
accuracy: 0.9746 - precision: 0.9792 - recall: 0.9706 - val_loss: 0.0919
- val_accuracy: 0.9724 - val_precision: 0.9765 - val_recall: 0.9706
Epoch 54/100
439/439 [=====] - 5s 12ms/step - loss: 0.0827 -
accuracy: 0.9737 - precision: 0.9790 - recall: 0.9706 - val_loss: 0.0897
- val_accuracy: 0.9736 - val_precision: 0.9774 - val_recall: 0.9716
Epoch 55/100
439/439 [=====] - 5s 12ms/step - loss: 0.0810 -
accuracy: 0.9750 - precision: 0.9791 - recall: 0.9705 - val_loss: 0.0894
- val_accuracy: 0.9732 - val_precision: 0.9768 - val_recall: 0.9716
Epoch 56/100
439/439 [=====] - 5s 12ms/step - loss: 0.0818 -
accuracy: 0.9751 - precision: 0.9793 - recall: 0.9714 - val_loss: 0.0887
- val_accuracy: 0.9736 - val_precision: 0.9769 - val_recall: 0.9711
Epoch 57/100
439/439 [=====] - 5s 12ms/step - loss: 0.0783 -
accuracy: 0.9754 - precision: 0.9790 - recall: 0.9723 - val_loss: 0.0880
- val_accuracy: 0.9734 - val_precision: 0.9772 - val_recall: 0.9713
Epoch 00057: early stopping
Test loss: 0.10033061355352402
Test accuracy: 0.9710000157356262
Test precision: 0.9738430380821228
Test recall: 0.9679999947547913
Test f1 score: 0.971
Test AUC for digit: 0 0.9930644081388708
Test AUC for digit: 1 0.9911542498484037
Test AUC for digit: 2 0.9863290741003511
Test AUC for digit: 3 0.9797077491850635
Test AUC for digit: 4 0.986564103792539
Test AUC for digit: 5 0.9774254759105543
Test AUC for digit: 6 0.9903174603174603
Test AUC for digit: 7 0.9857344204244699
Test AUC for digit: 8 0.9765873015873017
Test AUC for digit: 9 0.9711964862665278



Model with learning rate equals 0.00001 performed pretty well but it needed 54 epochs to achieve 97,1% accuracy (compared to 6 epochs using standard learning rate equals 0.001). This is because optimizer "jumped" too small distance searching best results, and it had to do many iterations to find the best weights.

Basic Multi-layer CNN

```
In [17]: model_cnn = keras.Sequential([
    layers.Conv2D(32, (3,3), activation="relu", input_shape=(28,28,1)),
    layers.MaxPooling2D (2,2),
    layers.Conv2D(64, (3,3), activation="relu"),
    layers.MaxPooling2D (2,2),

    layers.Flatten(),
    layers.Dropout (.5),
    layers.Dense(10, activation="softmax")
])
model_cnn.summary()
predict_model(model_cnn, [es], epochs=100)
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 26, 26, 32)	320

max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
flatten_5 (Flatten)	(None, 1600)	0
dropout_5 (Dropout)	(None, 1600)	0
dense_21 (Dense)	(None, 10)	16010

=====
 Total params: 34,826
 Trainable params: 34,826
 Non-trainable params: 0

Epoch 1/100

439/439 [=====] - 22s 28ms/step - loss: 0.7757 - accuracy: 0.7586 - precision: 0.8819 - recall: 0.6326 - val_loss: 0.0861 - val_accuracy: 0.9740 - val_precision: 0.9813 - val_recall: 0.9697

Epoch 2/100

439/439 [=====] - 3s 7ms/step - loss: 0.1093 - accuracy: 0.9654 - precision: 0.9719 - recall: 0.9608 - val_loss: 0.0568 - val_accuracy: 0.9833 - val_precision: 0.9858 - val_recall: 0.9804

Epoch 3/100

439/439 [=====] - 3s 7ms/step - loss: 0.0831 - accuracy: 0.9750 - precision: 0.9788 - recall: 0.9721 - val_loss: 0.0478 - val_accuracy: 0.9866 - val_precision: 0.9887 - val_recall: 0.9846

Epoch 4/100

439/439 [=====] - 3s 7ms/step - loss: 0.0692 - accuracy: 0.9781 - precision: 0.9811 - recall: 0.9756 - val_loss: 0.0391 - val_accuracy: 0.9886 - val_precision: 0.9910 - val_recall: 0.9869

Epoch 5/100

439/439 [=====] - 3s 7ms/step - loss: 0.0525 - accuracy: 0.9844 - precision: 0.9866 - recall: 0.9822 - val_loss: 0.0364 - val_accuracy: 0.9895 - val_precision: 0.9920 - val_recall: 0.9879

Epoch 6/100

439/439 [=====] - 3s 7ms/step - loss: 0.0548 - accuracy: 0.9820 - precision: 0.9843 - recall: 0.9804 - val_loss: 0.0351 - val_accuracy: 0.9899 - val_precision: 0.9918 - val_recall: 0.9890

Epoch 7/100

439/439 [=====] - 3s 7ms/step - loss: 0.0452 - accuracy: 0.9856 - precision: 0.9872 - recall: 0.9844 - val_loss: 0.0332 - val_accuracy: 0.9903 - val_precision: 0.9923 - val_recall: 0.9895

Epoch 8/100

439/439 [=====] - 3s 7ms/step - loss: 0.0416 - accuracy: 0.9865 - precision: 0.9884 - recall: 0.9850 - val_loss: 0.0349 - val_accuracy: 0.9898 - val_precision: 0.9910 - val_recall: 0.9892

Epoch 9/100

439/439 [=====] - 3s 7ms/step - loss: 0.0412 - accuracy: 0.9868 - precision: 0.9884 - recall: 0.9854 - val_loss: 0.0303 - val_accuracy: 0.9918 - val_precision: 0.9928 - val_recall: 0.9908

Epoch 10/100

439/439 [=====] - 3s 7ms/step - loss: 0.0366 - accuracy: 0.9881 - precision: 0.9893 - recall: 0.9870 - val_loss: 0.0307 - val_accuracy: 0.9911 - val_precision: 0.9919 - val_recall: 0.9899

Epoch 11/100

439/439 [=====] - 3s 7ms/step - loss: 0.0414 - accuracy: 0.9866 - precision: 0.9877 - recall: 0.9853 - val_loss: 0.0335 - val_accuracy: 0.9911 - val_precision: 0.9918 - val_recall: 0.9906

Epoch 12/100

439/439 [=====] - 3s 7ms/step - loss: 0.0345 - accuracy: 0.9890 - precision: 0.9897 - recall: 0.9881 - val_loss: 0.0316 - val_accuracy: 0.9911 - val_precision: 0.9922 - val_recall: 0.9903


Epoch 00012: early stopping

Test loss: 0.03780661150813103

Test accuracy: 0.9877142906188965

Test precision: 0.988834798336029

Test recall: 0.9868571162223816
 Test f1 score: 0.9877142857142858
 Test AUC for digit: 0 0.9986987490590519
 Test AUC for digit: 1 0.9946251332301316
 Test AUC for digit: 2 0.9912251490762208
 Test AUC for digit: 3 0.9914630429597434
 Test AUC for digit: 4 0.9935349833293058
 Test AUC for digit: 5 0.9949048474533845
 Test AUC for digit: 6 0.9943650793650793
 Test AUC for digit: 7 0.9923824507574553
 Test AUC for digit: 8 0.9917460317460318
 Test AUC for digit: 9 0.9890579103854441



Our first convolutional model with 2 convolutional layers was performing even better than fully connected neural networks. This model is not overfitted, because train and validation loss and accuracy are close to each other. It has only 34,826 parameters to train, so the training of such model is pretty fast. On test set model achieves 98.7% accuracy which is great result.

Different number of convolutional layers

```

In [18]: model_cnn_short = keras.Sequential([
            layers.Conv2D(32, (3,3), activation="relu", input_shape=(28,28,1)),
            layers.MaxPooling2D (2,2),

            layers.Flatten(),
            layers.Dropout (.5),
  
```

```

        layers.Dense(10, activation="softmax")
    ])
    model_cnn_short.summary()
    predict_model(model_cnn_short, [es], epochs=100)

```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_2 (MaxPooling2D)	(None, 13, 13, 32)	0
flatten_6 (Flatten)	(None, 5408)	0
dropout_6 (Dropout)	(None, 5408)	0
dense_22 (Dense)	(None, 10)	54090
Total params: 54,410		
Trainable params: 54,410		
Non-trainable params: 0		

Epoch 1/100

439/439 [=====] - 4s 8ms/step - loss: 0.7126 - accuracy: 0.7976 - precision: 0.9174 - recall: 0.6466 - val_loss: 0.1849 - val_accuracy: 0.9453 - val_precision: 0.9596 - val_recall: 0.9362

Epoch 2/100

439/439 [=====] - 3s 6ms/step - loss: 0.1896 - accuracy: 0.9442 - precision: 0.9575 - recall: 0.9324 - val_loss: 0.1285 - val_accuracy: 0.9629 - val_precision: 0.9724 - val_recall: 0.9567

Epoch 3/100

439/439 [=====] - 3s 6ms/step - loss: 0.1387 - accuracy: 0.9582 - precision: 0.9662 - recall: 0.9508 - val_loss: 0.0996 - val_accuracy: 0.9719 - val_precision: 0.9783 - val_recall: 0.9678

Epoch 4/100

439/439 [=====] - 3s 6ms/step - loss: 0.1196 - accuracy: 0.9648 - precision: 0.9723 - recall: 0.9583 - val_loss: 0.0876 - val_accuracy: 0.9737 - val_precision: 0.9791 - val_recall: 0.9713

Epoch 5/100

439/439 [=====] - 3s 6ms/step - loss: 0.1106 - accuracy: 0.9663 - precision: 0.9729 - recall: 0.9611 - val_loss: 0.0805 - val_accuracy: 0.9765 - val_precision: 0.9801 - val_recall: 0.9739

Epoch 6/100

439/439 [=====] - 3s 6ms/step - loss: 0.0965 - accuracy: 0.9702 - precision: 0.9750 - recall: 0.9659 - val_loss: 0.0732 - val_accuracy: 0.9788 - val_precision: 0.9829 - val_recall: 0.9759

Epoch 7/100

439/439 [=====] - 3s 6ms/step - loss: 0.0908 - accuracy: 0.9717 - precision: 0.9763 - recall: 0.9680 - val_loss: 0.0660 - val_accuracy: 0.9805 - val_precision: 0.9827 - val_recall: 0.9771

Epoch 8/100

439/439 [=====] - 3s 7ms/step - loss: 0.0875 - accuracy: 0.9728 - precision: 0.9768 - recall: 0.9696 - val_loss: 0.0633 - val_accuracy: 0.9824 - val_precision: 0.9852 - val_recall: 0.9799

Epoch 9/100

439/439 [=====] - 3s 7ms/step - loss: 0.0787 - accuracy: 0.9759 - precision: 0.9793 - recall: 0.9729 - val_loss: 0.0615 - val_accuracy: 0.9817 - val_precision: 0.9856 - val_recall: 0.9789

Epoch 10/100

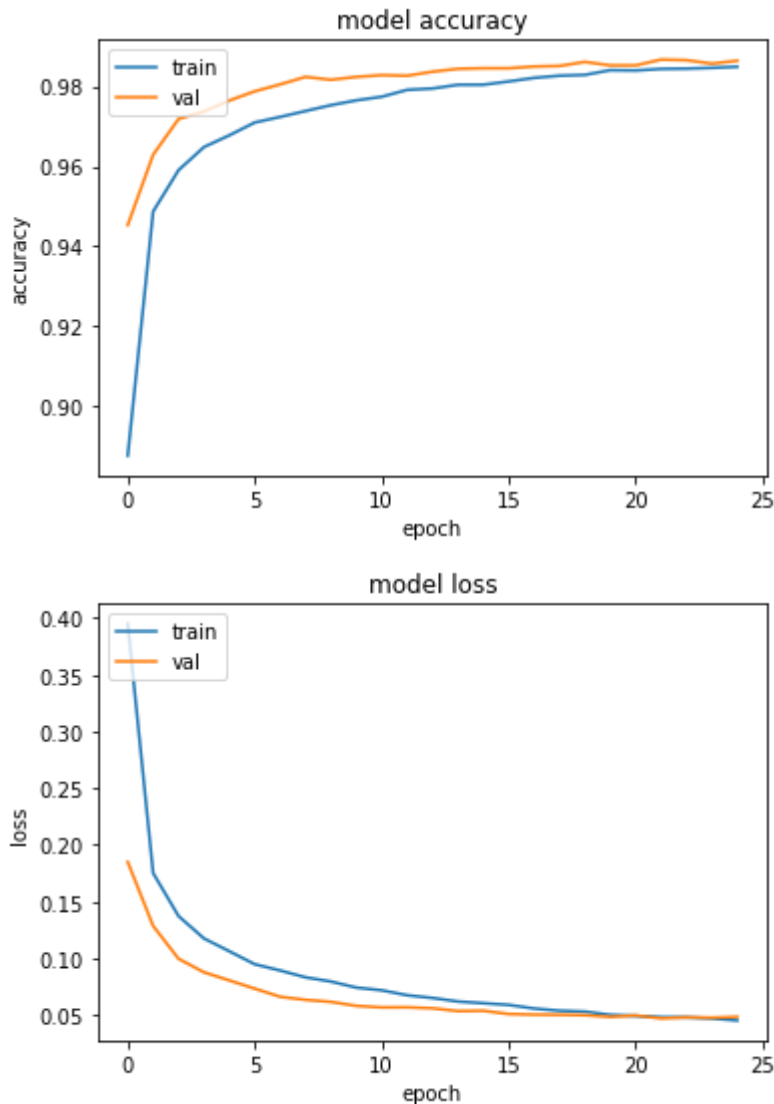
439/439 [=====] - 3s 6ms/step - loss: 0.0759 - accuracy: 0.9759 - precision: 0.9799 - recall: 0.9731 - val_loss: 0.0580 - val_accuracy: 0.9824 - val_precision: 0.9849 - val_recall: 0.9797

Epoch 11/100

439/439 [=====] - 3s 6ms/step - loss: 0.0730 - accuracy: 0.9771 - precision: 0.9806 - recall: 0.9748 - val_loss: 0.0567 - val_accuracy: 0.9828 - val_precision: 0.9849 - val_recall: 0.9798

Epoch 12/100
439/439 [=====] - 3s 6ms/step - loss: 0.0632 - accuracy: 0.9803 - precision: 0.9833 - recall: 0.9784 - val_loss: 0.0567 - val_accuracy: 0.9827 - val_precision: 0.9845 - val_recall: 0.9807
Epoch 13/100
439/439 [=====] - 3s 6ms/step - loss: 0.0649 - accuracy: 0.9793 - precision: 0.9820 - recall: 0.9770 - val_loss: 0.0559 - val_accuracy: 0.9837 - val_precision: 0.9854 - val_recall: 0.9820
Epoch 14/100
439/439 [=====] - 3s 6ms/step - loss: 0.0647 - accuracy: 0.9799 - precision: 0.9823 - recall: 0.9780 - val_loss: 0.0534 - val_accuracy: 0.9844 - val_precision: 0.9861 - val_recall: 0.9825
Epoch 15/100
439/439 [=====] - 3s 6ms/step - loss: 0.0569 - accuracy: 0.9815 - precision: 0.9842 - recall: 0.9795 - val_loss: 0.0537 - val_accuracy: 0.9846 - val_precision: 0.9857 - val_recall: 0.9834
Epoch 16/100
439/439 [=====] - 3s 6ms/step - loss: 0.0576 - accuracy: 0.9815 - precision: 0.9833 - recall: 0.9797 - val_loss: 0.0508 - val_accuracy: 0.9846 - val_precision: 0.9860 - val_recall: 0.9827
Epoch 17/100
439/439 [=====] - 3s 6ms/step - loss: 0.0528 - accuracy: 0.9825 - precision: 0.9848 - recall: 0.9807 - val_loss: 0.0503 - val_accuracy: 0.9850 - val_precision: 0.9864 - val_recall: 0.9843
Epoch 18/100
439/439 [=====] - 3s 6ms/step - loss: 0.0555 - accuracy: 0.9822 - precision: 0.9844 - recall: 0.9808 - val_loss: 0.0502 - val_accuracy: 0.9851 - val_precision: 0.9874 - val_recall: 0.9835
Epoch 19/100
439/439 [=====] - 3s 6ms/step - loss: 0.0516 - accuracy: 0.9831 - precision: 0.9852 - recall: 0.9814 - val_loss: 0.0498 - val_accuracy: 0.9861 - val_precision: 0.9877 - val_recall: 0.9848
Epoch 20/100
439/439 [=====] - 3s 6ms/step - loss: 0.0481 - accuracy: 0.9842 - precision: 0.9860 - recall: 0.9827 - val_loss: 0.0483 - val_accuracy: 0.9853 - val_precision: 0.9867 - val_recall: 0.9841
Epoch 21/100
439/439 [=====] - 3s 6ms/step - loss: 0.0466 - accuracy: 0.9856 - precision: 0.9873 - recall: 0.9842 - val_loss: 0.0496 - val_accuracy: 0.9853 - val_precision: 0.9870 - val_recall: 0.9843
Epoch 22/100
439/439 [=====] - 3s 6ms/step - loss: 0.0456 - accuracy: 0.9849 - precision: 0.9869 - recall: 0.9834 - val_loss: 0.0469 - val_accuracy: 0.9867 - val_precision: 0.9884 - val_recall: 0.9859
Epoch 23/100
439/439 [=====] - 3s 6ms/step - loss: 0.0454 - accuracy: 0.9850 - precision: 0.9869 - recall: 0.9837 - val_loss: 0.0479 - val_accuracy: 0.9866 - val_precision: 0.9876 - val_recall: 0.9848
Epoch 24/100
439/439 [=====] - 3s 6ms/step - loss: 0.0454 - accuracy: 0.9855 - precision: 0.9874 - recall: 0.9840 - val_loss: 0.0476 - val_accuracy: 0.9857 - val_precision: 0.9870 - val_recall: 0.9853
Epoch 25/100
439/439 [=====] - 3s 6ms/step - loss: 0.0442 - accuracy: 0.9851 - precision: 0.9868 - recall: 0.9840 - val_loss: 0.0483 - val_accuracy: 0.9864 - val_precision: 0.9871 - val_recall: 0.9853
Epoch 00025: early stopping
Test loss: 0.06184159591794014
Test accuracy: 0.9821428656578064
Test precision: 0.9840951561927795
Test recall: 0.9811428785324097
Test f1 score: 0.9821428571428571
Test AUC for digit: 0 0.9946403781193142
Test AUC for digit: 1 0.9928970988504665
Test AUC for digit: 2 0.9918050336696427
Test AUC for digit: 3 0.9915450335250312
Test AUC for digit: 4 0.9882236580836722
Test AUC for digit: 5 0.9855790503765007

Test AUC for digit: 6 0.9932539682539682
 Test AUC for digit: 7 0.9851409580367276
 Test AUC for digit: 8 0.9859523809523809
 Test AUC for digit: 9 0.9908959987735204



Next model has only 1 convolution layer which has more parameters (54,410) because of less number of pooling layers. The results are satisfying, but not as good as previous model (test accuracy equals 98,2%).

```
In [67]: model_cnn_long = keras.Sequential([
    layers.Conv2D(32, (3,3), activation="relu", input_shape=(28,28,1)),
    layers.MaxPooling2D ((2,2),1),
    layers.Conv2D(64, (3,3), activation="relu"),
    layers.MaxPooling2D ((2,2),1),
    layers.Conv2D(128, (3,3), activation="relu"),
    layers.MaxPooling2D ((2,2),1),
    layers.Conv2D(512, (3,3), activation="relu"),
    layers.MaxPooling2D ((2,2),1),

    layers.Flatten(),
    layers.Dense(128, activation="relu"),
    layers.Dropout(.5),
    layers.Dense(10, activation="softmax")
])
model_cnn_long.summary()
predict_model(model_cnn_long, [es], epochs=100)
```

Model: "sequential_28"

Layer (type)	Output Shape	Param #
conv2d_32 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_17 (MaxPooling)	(None, 25, 25, 32)	0
conv2d_33 (Conv2D)	(None, 23, 23, 64)	18496
max_pooling2d_18 (MaxPooling)	(None, 22, 22, 64)	0
conv2d_34 (Conv2D)	(None, 20, 20, 128)	73856
max_pooling2d_19 (MaxPooling)	(None, 19, 19, 128)	0
conv2d_35 (Conv2D)	(None, 17, 17, 512)	590336
max_pooling2d_20 (MaxPooling)	(None, 16, 16, 512)	0
flatten_28 (Flatten)	(None, 131072)	0
dense_85 (Dense)	(None, 128)	16777344
dropout_22 (Dropout)	(None, 128)	0
dense_86 (Dense)	(None, 10)	1290
Total params: 17,461,642		
Trainable params: 17,461,642		
Non-trainable params: 0		

Epoch 1/100

439/439 [=====] - 19s 42ms/step - loss: 0.5369 - accuracy: 0.8446 - precision: 0.9042 - recall: 0.7958 - val_loss: 0.0644 - val_accuracy: 0.9812 - val_precision: 0.9828 - val_recall: 0.9805

Epoch 2/100

439/439 [=====] - 18s 40ms/step - loss: 0.0900 - accuracy: 0.9740 - precision: 0.9778 - recall: 0.9702 - val_loss: 0.0481 - val_accuracy: 0.9846 - val_precision: 0.9861 - val_recall: 0.9833

Epoch 3/100

439/439 [=====] - 18s 40ms/step - loss: 0.0727 - accuracy: 0.9790 - precision: 0.9818 - recall: 0.9766 - val_loss: 0.0360 - val_accuracy: 0.9880 - val_precision: 0.9900 - val_recall: 0.9864

Epoch 4/100

439/439 [=====] - 18s 41ms/step - loss: 0.0553 - accuracy: 0.9824 - precision: 0.9848 - recall: 0.9804 - val_loss: 0.0303 - val_accuracy: 0.9915 - val_precision: 0.9922 - val_recall: 0.9909

Epoch 5/100

439/439 [=====] - 18s 40ms/step - loss: 0.0370 - accuracy: 0.9888 - precision: 0.9900 - recall: 0.9874 - val_loss: 0.0300 - val_accuracy: 0.9925 - val_precision: 0.9935 - val_recall: 0.9922

Epoch 6/100

439/439 [=====] - 18s 40ms/step - loss: 0.0322 - accuracy: 0.9895 - precision: 0.9908 - recall: 0.9885 - val_loss: 0.0272 - val_accuracy: 0.9941 - val_precision: 0.9942 - val_recall: 0.9935

Epoch 7/100

439/439 [=====] - 18s 41ms/step - loss: 0.0295 - accuracy: 0.9911 - precision: 0.9922 - recall: 0.9904 - val_loss: 0.0321 - val_accuracy: 0.9916 - val_precision: 0.9923 - val_recall: 0.9915

Epoch 8/100

439/439 [=====] - 18s 41ms/step - loss: 0.0277 - accuracy: 0.9911 - precision: 0.9924 - recall: 0.9905 - val_loss: 0.0221 - val_accuracy: 0.9932 - val_precision: 0.9941 - val_recall: 0.9931


Epoch 9/100

439/439 [=====] - 18s 41ms/step - loss: 0.0235 - accuracy: 0.9925 - precision: 0.9931 - recall: 0.9920 - val_loss: 0.0280 - val_accuracy: 0.9924 - val_precision: 0.9925 - val_recall: 0.9924

```

Restoring model weights from the end of the best epoch.
Epoch 00009: early stopping
Test loss: 0.03389483317732811
Test accuracy: 0.9918571710586548
Test precision: 0.992139458656311
Test recall: 0.9917142987251282
Test f1 score: 0.9918571428571429
Test AUC for digit: 0 0.9990376444107582
Test AUC for digit: 1 0.9981951123812957
Test AUC for digit: 2 0.9929911952983329
Test AUC for digit: 3 0.9945313397184135
Test AUC for digit: 4 0.9947926711881807
Test AUC for digit: 5 0.9956330341254677
Test AUC for digit: 6 0.9968966508108128
Test AUC for digit: 7 0.999361124421019
Test AUC for digit: 8 0.9925846524567535
Test AUC for digit: 9 0.990134076231021

```



Next we created a neural network with 4 convolutional layers and with 17 million parameters. The model was not overfitted. It had accuracy around 99.2% for test, train and validation model. Time needed to train this model was much higher (19s per epoch comparing to 3s per epoch in CNN that we have implemented). This is the best model that we have created for this dataset.

Different number of filters per layer

```
In [21]: model_cnn_min = keras.Sequential([
```

```

layers.Conv2D(4, (3,3), activation="relu", input_shape=(28,28,1)),
layers.MaxPooling2D (2,2),
layers.Conv2D(16, (3,3), activation="relu"),
layers.MaxPooling2D (2,2),

layers.Flatten(),
layers.Dropout(.5),
layers.Dense(10, activation="softmax")
])
model_cnn_min.summary()
predict_model(model_cnn_min, [es], epochs=100)

```

Model: "sequential_9"

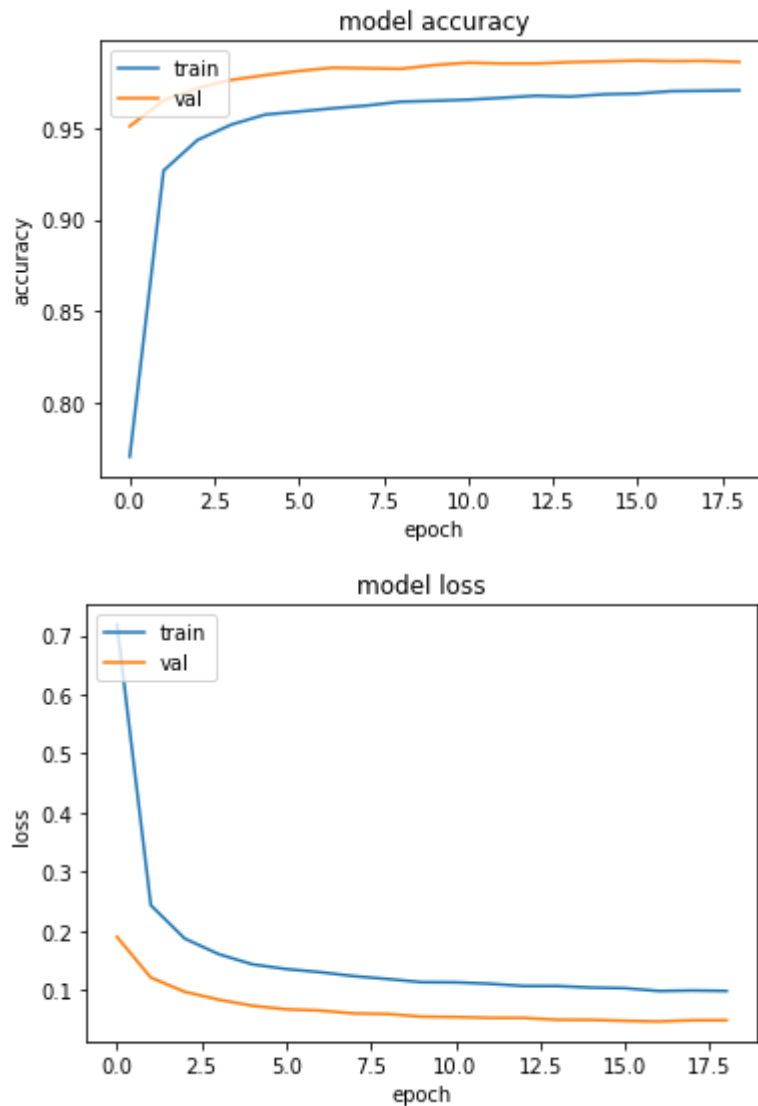
Layer (type)	Output Shape	Param #
conv2d_9 (Conv2D)	(None, 26, 26, 4)	40
max_pooling2d_7 (MaxPooling2D)	(None, 13, 13, 4)	0
conv2d_10 (Conv2D)	(None, 11, 11, 16)	592
max_pooling2d_8 (MaxPooling2D)	(None, 5, 5, 16)	0
flatten_9 (Flatten)	(None, 400)	0
dropout_8 (Dropout)	(None, 400)	0
dense_28 (Dense)	(None, 10)	4010
Total params: 4,642		
Trainable params: 4,642		
Non-trainable params: 0		

```

Epoch 1/100
439/439 [=====] - 20s 29ms/step - loss: 1.2412 - accuracy: 0.5863 - precision: 0.8270 - recall: 0.3904 - val_loss: 0.1895 - val_accuracy: 0.9508 - val_precision: 0.9644 - val_recall: 0.9309
Epoch 2/100
439/439 [=====] - 3s 6ms/step - loss: 0.2660 - accuracy: 0.9202 - precision: 0.9399 - recall: 0.9013 - val_loss: 0.1206 - val_accuracy: 0.9649 - val_precision: 0.9742 - val_recall: 0.9571
Epoch 3/100
439/439 [=====] - 3s 6ms/step - loss: 0.1972 - accuracy: 0.9398 - precision: 0.9519 - recall: 0.9284 - val_loss: 0.0964 - val_accuracy: 0.9716 - val_precision: 0.9768 - val_recall: 0.9651
Epoch 4/100
439/439 [=====] - 3s 6ms/step - loss: 0.1631 - accuracy: 0.9512 - precision: 0.9599 - recall: 0.9423 - val_loss: 0.0831 - val_accuracy: 0.9760 - val_precision: 0.9807 - val_recall: 0.9701
Epoch 5/100
439/439 [=====] - 3s 6ms/step - loss: 0.1448 - accuracy: 0.9565 - precision: 0.9640 - recall: 0.9497 - val_loss: 0.0728 - val_accuracy: 0.9786 - val_precision: 0.9820 - val_recall: 0.9747
Epoch 6/100
439/439 [=====] - 3s 6ms/step - loss: 0.1358 - accuracy: 0.9586 - precision: 0.9660 - recall: 0.9528 - val_loss: 0.0668 - val_accuracy: 0.9810 - val_precision: 0.9846 - val_recall: 0.9784
Epoch 7/100
439/439 [=====] - 3s 6ms/step - loss: 0.1321 - accuracy: 0.9604 - precision: 0.9674 - recall: 0.9548 - val_loss: 0.0647 - val_accuracy: 0.9827 - val_precision: 0.9860 - val_recall: 0.9788
Epoch 8/100
439/439 [=====] - 3s 6ms/step - loss: 0.1219 - accuracy: 0.9629 - precision: 0.9683 - recall: 0.9573 - val_loss: 0.0597 - val_accuracy: 0.9824 - val_precision: 0.9862 - val_recall: 0.9801
Epoch 9/100

```

439/439 [=====] - 3s 6ms/step - loss: 0.1195 - accuracy: 0.9628 - precision: 0.9687 - recall: 0.9579 - val_loss: 0.0588 - val_accuracy: 0.9821 - val_precision: 0.9861 - val_recall: 0.9805
Epoch 10/100
439/439 [=====] - 3s 6ms/step - loss: 0.1119 - accuracy: 0.9654 - precision: 0.9695 - recall: 0.9610 - val_loss: 0.0542 - val_accuracy: 0.9841 - val_precision: 0.9865 - val_recall: 0.9818
Epoch 11/100
439/439 [=====] - 3s 6ms/step - loss: 0.1154 - accuracy: 0.9644 - precision: 0.9700 - recall: 0.9598 - val_loss: 0.0537 - val_accuracy: 0.9854 - val_precision: 0.9875 - val_recall: 0.9827
Epoch 12/100
439/439 [=====] - 3s 6ms/step - loss: 0.1106 - accuracy: 0.9667 - precision: 0.9717 - recall: 0.9622 - val_loss: 0.0522 - val_accuracy: 0.9850 - val_precision: 0.9875 - val_recall: 0.9827
Epoch 13/100
439/439 [=====] - 3s 6ms/step - loss: 0.1048 - accuracy: 0.9677 - precision: 0.9727 - recall: 0.9641 - val_loss: 0.0522 - val_accuracy: 0.9850 - val_precision: 0.9868 - val_recall: 0.9830
Epoch 14/100
439/439 [=====] - 3s 6ms/step - loss: 0.1069 - accuracy: 0.9672 - precision: 0.9715 - recall: 0.9638 - val_loss: 0.0490 - val_accuracy: 0.9857 - val_precision: 0.9877 - val_recall: 0.9843
Epoch 15/100
439/439 [=====] - 3s 6ms/step - loss: 0.1051 - accuracy: 0.9676 - precision: 0.9712 - recall: 0.9637 - val_loss: 0.0488 - val_accuracy: 0.9861 - val_precision: 0.9875 - val_recall: 0.9835
Epoch 16/100
439/439 [=====] - 3s 6ms/step - loss: 0.1029 - accuracy: 0.9683 - precision: 0.9719 - recall: 0.9650 - val_loss: 0.0472 - val_accuracy: 0.9866 - val_precision: 0.9884 - val_recall: 0.9851
Epoch 17/100
439/439 [=====] - 3s 6ms/step - loss: 0.0990 - accuracy: 0.9694 - precision: 0.9730 - recall: 0.9657 - val_loss: 0.0461 - val_accuracy: 0.9863 - val_precision: 0.9883 - val_recall: 0.9847
Epoch 18/100
439/439 [=====] - 3s 6ms/step - loss: 0.0960 - accuracy: 0.9705 - precision: 0.9743 - recall: 0.9670 - val_loss: 0.0482 - val_accuracy: 0.9864 - val_precision: 0.9881 - val_recall: 0.9838
Epoch 19/100
439/439 [=====] - 3s 6ms/step - loss: 0.0978 - accuracy: 0.9701 - precision: 0.9744 - recall: 0.9665 - val_loss: 0.0482 - val_accuracy: 0.9859 - val_precision: 0.9877 - val_recall: 0.9840
Epoch 00019: early stopping
Test loss: 0.06978869438171387
Test accuracy: 0.9787142872810364
Test precision: 0.9813513159751892
Test recall: 0.9772857427597046
Test f1 score: 0.9787142857142858
Test AUC for digit: 0 0.9965322040694353
Test AUC for digit: 1 0.9950131545729289
Test AUC for digit: 2 0.9902744281128235
Test AUC for digit: 3 0.9815570919354413
Test AUC for digit: 4 0.9918910464386338
Test AUC for digit: 5 0.9894167367175205
Test AUC for digit: 6 0.9894444444444445
Test AUC for digit: 7 0.9815003368863819
Test AUC for digit: 8 0.9804761904761905
Test AUC for digit: 9 0.9860031281752693



Next we decided to check how number of filters impact to performance of model. Reducing number of filter in convolutional layers made our model worse than basic model. Accuracy has fallen to 97.8%, because this model was too simple to explain complexity of our data. This model is underfitted.

```
In [22]: model_cnn_max = keras.Sequential([
    layers.Conv2D(128, (3,3), activation="relu", input_shape=(28,28,1)),
    layers.MaxPooling2D (2,2),
    layers.Conv2D(512, (3,3), activation="relu"),
    layers.MaxPooling2D (2,2),

    layers.Flatten(),
    layers.Dropout(.5),
    layers.Dense(10, activation="softmax")
])
model_cnn_max.summary()
predict_model(model_cnn_max, [es], epochs=100)
```

Model: "sequential_10"

Layer (type)	Output Shape	Param #
=====		
conv2d_11 (Conv2D)	(None, 26, 26, 128)	1280
=====		
max_pooling2d_9 (MaxPooling2D)	(None, 13, 13, 128)	0
=====		

conv2d_12 (Conv2D)	(None, 11, 11, 512)	590336
max_pooling2d_10 (MaxPooling)	(None, 5, 5, 512)	0
flatten_10 (Flatten)	(None, 12800)	0
dropout_9 (Dropout)	(None, 12800)	0
dense_29 (Dense)	(None, 10)	128010


=====
 Total params: 719,626
 Trainable params: 719,626
 Non-trainable params: 0

Epoch 1/100	439/439 [=====] - 27s 44ms/step - loss: 0.3990 - accuracy: 0.8824 - precision: 0.9333 - recall: 0.8202 - val_loss: 0.0549 - val_accuracy: 0.9824 - val_precision: 0.9854 - val_recall: 0.9804
Epoch 2/100	439/439 [=====] - 9s 19ms/step - loss: 0.0521 - accuracy: 0.9837 - precision: 0.9856 - recall: 0.9818 - val_loss: 0.0435 - val_accuracy: 0.9857 - val_precision: 0.9870 - val_recall: 0.9844
Epoch 3/100	439/439 [=====] - 8s 19ms/step - loss: 0.0429 - accuracy: 0.9862 - precision: 0.9876 - recall: 0.9849 - val_loss: 0.0343 - val_accuracy: 0.9896 - val_precision: 0.9902 - val_recall: 0.9885
Epoch 4/100	439/439 [=====] - 8s 19ms/step - loss: 0.0320 - accuracy: 0.9897 - precision: 0.9904 - recall: 0.9890 - val_loss: 0.0361 - val_accuracy: 0.9899 - val_precision: 0.9909 - val_recall: 0.9895
Epoch 5/100	439/439 [=====] - 8s 19ms/step - loss: 0.0249 - accuracy: 0.9921 - precision: 0.9929 - recall: 0.9916 - val_loss: 0.0323 - val_accuracy: 0.9900 - val_precision: 0.9913 - val_recall: 0.9890
Epoch 6/100	439/439 [=====] - 8s 19ms/step - loss: 0.0233 - accuracy: 0.9925 - precision: 0.9928 - recall: 0.9920 - val_loss: 0.0347 - val_accuracy: 0.9899 - val_precision: 0.9903 - val_recall: 0.9896
Epoch 7/100	439/439 [=====] - 8s 19ms/step - loss: 0.0185 - accuracy: 0.9941 - precision: 0.9945 - recall: 0.9938 - val_loss: 0.0343 - val_accuracy: 0.9915 - val_precision: 0.9921 - val_recall: 0.9911
Epoch 8/100	439/439 [=====] - 8s 19ms/step - loss: 0.0156 - accuracy: 0.9946 - precision: 0.9950 - recall: 0.9943 - val_loss: 0.0346 - val_accuracy: 0.9906 - val_precision: 0.9908 - val_recall: 0.9903
Epoch 9/100	439/439 [=====] - 8s 19ms/step - loss: 0.0137 - accuracy: 0.9956 - precision: 0.9960 - recall: 0.9954 - val_loss: 0.0362 - val_accuracy: 0.9899 - val_precision: 0.9906 - val_recall: 0.9899
Epoch 10/100	439/439 [=====] - 8s 19ms/step - loss: 0.0112 - accuracy: 0.9961 - precision: 0.9966 - recall: 0.9959 - val_loss: 0.0365 - val_accuracy: 0.9903 - val_precision: 0.9909 - val_recall: 0.9902

Epoch 00010: early stopping
 Test loss: 0.04324812442064285
 Test accuracy: 0.9902856945991516
 Test precision: 0.9904244542121887
 Test recall: 0.9900000095367432
 Test f1 score: 0.9902857142857143
 Test AUC for digit: 0 0.9962574850299402
 Test AUC for digit: 1 0.9940907921236538
 Test AUC for digit: 2 0.9984028163786178
 Test AUC for digit: 3 0.993960028357141
 Test AUC for digit: 4 0.9914325054702
 Test AUC for digit: 5 0.9978038687780799
 Test AUC for digit: 6 0.9957936507936508
 Test AUC for digit: 7 0.9913140793939336

Test AUC for digit: 8 0.9926190476190476

Test AUC for digit: 9 0.9948151881227197



Next we increased number of filters. This caused a raise of number of parameters to over 700 thousands but model did not perform better than basic model. Test accuracy was 99%, which is slightly less than basic model's accuracy. It means that we should not use such high number of filters because we do not need them. This model also seems to be overfitted.

Different size and type of pooling layers

```
In [41]: model_cnn_pool5 = keras.Sequential([
    layers.Conv2D(32, (3,3), activation="relu", input_shape=(28,28,1)),
    layers.MaxPooling2D (5,3),
    layers.Conv2D(64, (3,3), activation="relu"),
    layers.MaxPooling2D (5,3),

    layers.Flatten(),
    layers.Dropout(.5),
    layers.Dense(10, activation="softmax")
])
model_cnn_pool5.summary()
predict_model(model_cnn_pool5, [es], epochs=100)
```

Model: "sequential_17"

Layer (type)	Output Shape	Param #
conv2d_13 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_11 (MaxPooling)	(None, 8, 8, 32)	0
conv2d_14 (Conv2D)	(None, 6, 6, 64)	18496
max_pooling2d_12 (MaxPooling)	(None, 1, 1, 64)	0
flatten_17 (Flatten)	(None, 64)	0
dropout_16 (Dropout)	(None, 64)	0
dense_54 (Dense)	(None, 10)	650

Total params: 19,466
 Trainable params: 19,466
 Non-trainable params: 0

Epoch 1/100
 439/439 [=====] - 17s 25ms/step - loss: 1.6386 - accuracy: 0.4443 - precision: 0.7621 - recall: 0.1895 - val_loss: 0.3402 - val_accuracy: 0.9205 - val_precision: 0.9659 - val_recall: 0.8592

Epoch 2/100
 439/439 [=====] - 3s 8ms/step - loss: 0.5372 - accuracy: 0.8335 - precision: 0.8979 - recall: 0.7583 - val_loss: 0.2031 - val_accuracy: 0.9482 - val_precision: 0.9721 - val_recall: 0.9242

Epoch 3/100
 439/439 [=====] - 3s 8ms/step - loss: 0.3899 - accuracy: 0.8815 - precision: 0.9224 - recall: 0.8415 - val_loss: 0.1599 - val_accuracy: 0.9600 - val_precision: 0.9747 - val_recall: 0.9439

Epoch 4/100
 439/439 [=====] - 3s 7ms/step - loss: 0.3313 - accuracy: 0.8988 - precision: 0.9294 - recall: 0.8673 - val_loss: 0.1329 - val_accuracy: 0.9629 - val_precision: 0.9762 - val_recall: 0.9489

Epoch 5/100
 439/439 [=====] - 3s 7ms/step - loss: 0.2954 - accuracy: 0.9092 - precision: 0.9359 - recall: 0.8834 - val_loss: 0.1220 - val_accuracy: 0.9681 - val_precision: 0.9795 - val_recall: 0.9576

Epoch 6/100
 439/439 [=====] - 3s 8ms/step - loss: 0.2813 - accuracy: 0.9143 - precision: 0.9392 - recall: 0.8897 - val_loss: 0.1032 - val_accuracy: 0.9701 - val_precision: 0.9799 - val_recall: 0.9633

Epoch 7/100
 439/439 [=====] - 3s 8ms/step - loss: 0.2496 - accuracy: 0.9239 - precision: 0.9440 - recall: 0.9045 - val_loss: 0.1051 - val_accuracy: 0.9704 - val_precision: 0.9797 - val_recall: 0.9593

Epoch 8/100
 439/439 [=====] - 3s 7ms/step - loss: 0.2431 - accuracy: 0.9248 - precision: 0.9447 - recall: 0.9058 - val_loss: 0.0936 - val_accuracy: 0.9706 - val_precision: 0.9798 - val_recall: 0.9641

Epoch 9/100
 439/439 [=====] - 3s 7ms/step - loss: 0.2384 - accuracy: 0.9260 - precision: 0.9459 - recall: 0.9090 - val_loss: 0.0853 - val_accuracy: 0.9766 - val_precision: 0.9830 - val_recall: 0.9696

Epoch 10/100
 439/439 [=====] - 3s 7ms/step - loss: 0.2130 - accuracy: 0.9333 - precision: 0.9492 - recall: 0.9185 - val_loss: 0.0826 - val_accuracy: 0.9762 - val_precision: 0.9835 - val_recall: 0.9711

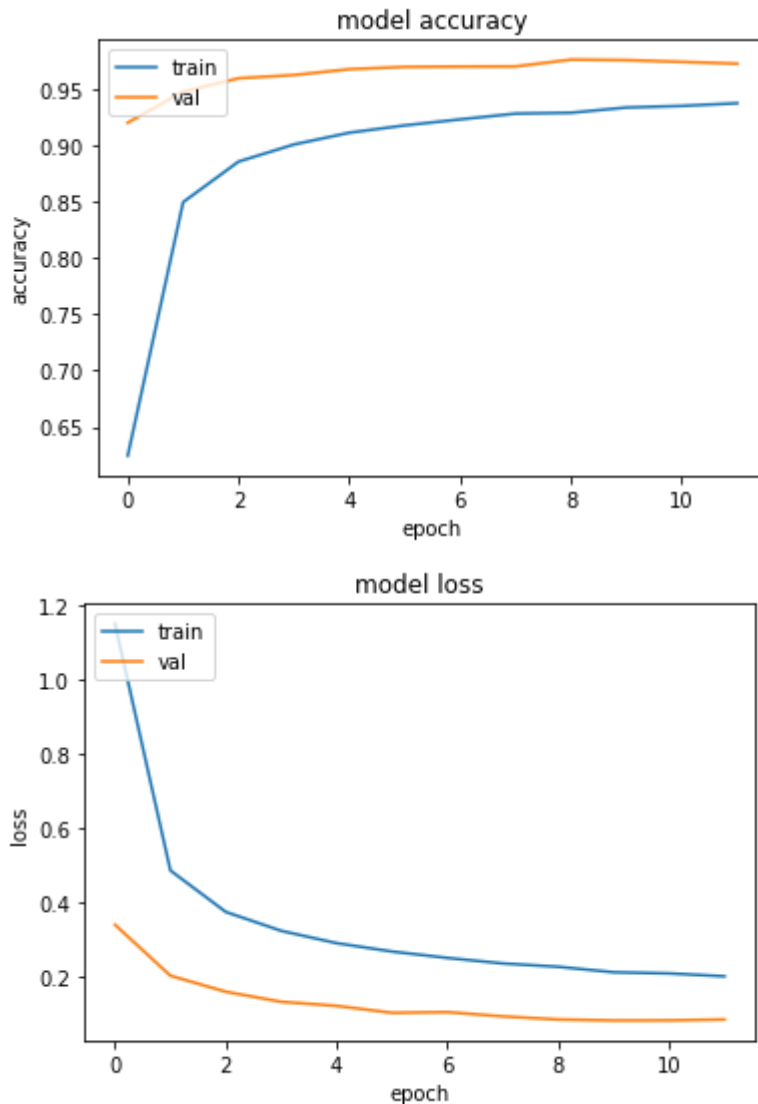
Epoch 11/100
 439/439 [=====] - 3s 7ms/step - loss: 0.2164 - accuracy: 0.9316 - precision: 0.9492 - recall: 0.9175 - val_loss: 0.0828 - val_accuracy: 0.9746 - val_precision: 0.9807 - val_recall: 0.9701

Epoch 12/100
 439/439 [=====] - 3s 7ms/step - loss: 0.2055 - accuracy: 0.9370 - precision: 0.9515 - recall: 0.9241 - val_loss: 0.0853 - val_accuracy: 0.9730 - val_precision: 0.9796 - val_recall: 0.9698

```

Epoch 00012: early stopping
Test loss: 0.09769929945468903
Test accuracy: 0.9715714454650879
Test precision: 0.9770860075950623
Test recall: 0.9685714244842529
Test f1 score: 0.9715714285714285
Test AUC for digit: 0 0.9948361331663899
Test AUC for digit: 1 0.9941565446942954
Test AUC for digit: 2 0.9841899519327072
Test AUC for digit: 3 0.9719724776720643
Test AUC for digit: 4 0.9860265364292893
Test AUC for digit: 5 0.9800909396406366
Test AUC for digit: 6 0.9928571428571429
Test AUC for digit: 7 0.9823115528040143
Test AUC for digit: 8 0.9796031746031746
Test AUC for digit: 9 0.9761439939197929

```



Next, we checked different size of pooling layers. We decided to create a MaxPooling layers with size equals (5,5) and stride equals 3. It means that we take a square of values with size 5x5 then we look for max value, we write it in the middle of square and than we "move" 3 numbers in right or down direction. As we can see, the accuracy is worse than basic model, because we lose too much information in MaxPooling layers. The plots also shows that this model is underfitted.

```

In [42]: model_cnn_avg = keras.Sequential([
          layers.Conv2D(32, (3,3), activation="relu", input_shape=(28,28,1)),
          layers.AveragePooling2D (3,3),

```

```

layers.Conv2D(64, (3,3), activation="relu"),
layers.AveragePooling2D (3,3),

layers.Flatten(),
layers.Dropout(.5),
layers.Dense(10, activation="softmax")
])
model_cnn_avg.summary()
predict_model(model_cnn_avg, [es], epochs=100)

```

Model: "sequential_18"

Layer (type)	Output Shape	Param #
conv2d_15 (Conv2D)	(None, 26, 26, 32)	320
average_pooling2d_2 (Average)	(None, 8, 8, 32)	0
conv2d_16 (Conv2D)	(None, 6, 6, 64)	18496
average_pooling2d_3 (Average)	(None, 2, 2, 64)	0
flatten_18 (Flatten)	(None, 256)	0
dropout_17 (Dropout)	(None, 256)	0
dense_55 (Dense)	(None, 10)	2570
Total params: 21,386		
Trainable params: 21,386		
Non-trainable params: 0		

Epoch 1/100

439/439 [=====] - 4s 8ms/step - loss: 1.4358 - accuracy: 0.5197 - precision: 0.7783 - recall: 0.3063 - val_loss: 0.2705 - val_accuracy: 0.9248 - val_precision: 0.9531 - val_recall: 0.8922

Epoch 2/100

439/439 [=====] - 3s 8ms/step - loss: 0.3783 - accuracy: 0.8837 - precision: 0.9196 - recall: 0.8476 - val_loss: 0.1960 - val_accuracy: 0.9437 - val_precision: 0.9590 - val_recall: 0.9251

Epoch 3/100

439/439 [=====] - 3s 7ms/step - loss: 0.2821 - accuracy: 0.9156 - precision: 0.9367 - recall: 0.8952 - val_loss: 0.1518 - val_accuracy: 0.9548 - val_precision: 0.9683 - val_recall: 0.9423

Epoch 4/100

439/439 [=====] - 3s 7ms/step - loss: 0.2438 - accuracy: 0.9272 - precision: 0.9439 - recall: 0.9095 - val_loss: 0.1290 - val_accuracy: 0.9618 - val_precision: 0.9702 - val_recall: 0.9535

Epoch 5/100

439/439 [=====] - 3s 7ms/step - loss: 0.2236 - accuracy: 0.9325 - precision: 0.9481 - recall: 0.9185 - val_loss: 0.1161 - val_accuracy: 0.9654 - val_precision: 0.9729 - val_recall: 0.9576

Epoch 6/100

439/439 [=====] - 3s 7ms/step - loss: 0.2054 - accuracy: 0.9364 - precision: 0.9496 - recall: 0.9246 - val_loss: 0.1074 - val_accuracy: 0.9674 - val_precision: 0.9733 - val_recall: 0.9622

Epoch 7/100

439/439 [=====] - 3s 7ms/step - loss: 0.1861 - accuracy: 0.9435 - precision: 0.9551 - recall: 0.9341 - val_loss: 0.1005 - val_accuracy: 0.9729 - val_precision: 0.9769 - val_recall: 0.9655

Epoch 8/100


439/439 [=====] - 3s 7ms/step - loss: 0.1737 - accuracy: 0.9474 - precision: 0.9569 - recall: 0.9373 - val_loss: 0.0930 - val_accuracy: 0.9732 - val_precision: 0.9797 - val_recall: 0.9664


Epoch 9/100

439/439 [=====] - 3s 7ms/step - loss: 0.1644 - accuracy: 0.9495 - precision: 0.9593 - recall: 0.9412 - val_loss: 0.0846 -

val_accuracy: 0.9769 - val_precision: 0.9813 - val_recall: 0.9713
Epoch 10/100
439/439 [=====] - 3s 7ms/step - loss: 0.1554 - a
ccuracy: 0.9541 - precision: 0.9624 - recall: 0.9459 - val_loss: 0.0813 -
val_accuracy: 0.9769 - val_precision: 0.9796 - val_recall: 0.9720
Epoch 11/100
439/439 [=====] - 3s 7ms/step - loss: 0.1508 - a
ccuracy: 0.9546 - precision: 0.9632 - recall: 0.9462 - val_loss: 0.0804 -
val_accuracy: 0.9788 - val_precision: 0.9818 - val_recall: 0.9733
Epoch 12/100
439/439 [=====] - 3s 7ms/step - loss: 0.1422 - a
ccuracy: 0.9573 - precision: 0.9656 - recall: 0.9500 - val_loss: 0.0769 -
val_accuracy: 0.9784 - val_precision: 0.9820 - val_recall: 0.9747
Epoch 13/100
439/439 [=====] - 3s 7ms/step - loss: 0.1312 - a
ccuracy: 0.9601 - precision: 0.9681 - recall: 0.9534 - val_loss: 0.0703 -
val_accuracy: 0.9805 - val_precision: 0.9837 - val_recall: 0.9772
Epoch 14/100
439/439 [=====] - 3s 7ms/step - loss: 0.1299 - a
ccuracy: 0.9603 - precision: 0.9674 - recall: 0.9546 - val_loss: 0.0670 -
val_accuracy: 0.9812 - val_precision: 0.9850 - val_recall: 0.9778
Epoch 15/100
439/439 [=====] - 3s 7ms/step - loss: 0.1196 - a
ccuracy: 0.9636 - precision: 0.9695 - recall: 0.9576 - val_loss: 0.0651 -
val_accuracy: 0.9814 - val_precision: 0.9845 - val_recall: 0.9785
Epoch 16/100
439/439 [=====] - 3s 7ms/step - loss: 0.1184 - a
ccuracy: 0.9636 - precision: 0.9692 - recall: 0.9584 - val_loss: 0.0651 -
val_accuracy: 0.9821 - val_precision: 0.9855 - val_recall: 0.9779
Epoch 17/100
439/439 [=====] - 3s 7ms/step - loss: 0.1094 - a
ccuracy: 0.9654 - precision: 0.9711 - recall: 0.9599 - val_loss: 0.0589 -
val_accuracy: 0.9834 - val_precision: 0.9855 - val_recall: 0.9808
Epoch 18/100
439/439 [=====] - 3s 7ms/step - loss: 0.1074 - a
ccuracy: 0.9668 - precision: 0.9724 - recall: 0.9627 - val_loss: 0.0578 -
val_accuracy: 0.9837 - val_precision: 0.9866 - val_recall: 0.9811
Epoch 19/100
439/439 [=====] - 3s 7ms/step - loss: 0.1010 - a
ccuracy: 0.9683 - precision: 0.9735 - recall: 0.9638 - val_loss: 0.0559 -
val_accuracy: 0.9843 - val_precision: 0.9865 - val_recall: 0.9817
Epoch 20/100
439/439 [=====] - 3s 7ms/step - loss: 0.0975 - a
ccuracy: 0.9704 - precision: 0.9751 - recall: 0.9660 - val_loss: 0.0564 -
val_accuracy: 0.9833 - val_precision: 0.9867 - val_recall: 0.9817
Epoch 21/100
439/439 [=====] - 3s 7ms/step - loss: 0.0961 - a
ccuracy: 0.9713 - precision: 0.9754 - recall: 0.9670 - val_loss: 0.0562 -
val_accuracy: 0.9846 - val_precision: 0.9865 - val_recall: 0.9821
Epoch 22/100
439/439 [=====] - 3s 7ms/step - loss: 0.0930 - a
ccuracy: 0.9708 - precision: 0.9752 - recall: 0.9669 - val_loss: 0.0506 -
val_accuracy: 0.9856 - val_precision: 0.9887 - val_recall: 0.9843
Epoch 23/100
439/439 [=====] - 3s 7ms/step - loss: 0.0942 - a
ccuracy: 0.9710 - precision: 0.9755 - recall: 0.9667 - val_loss: 0.0512 -
val_accuracy: 0.9843 - val_precision: 0.9869 - val_recall: 0.9821
Epoch 24/100
439/439 [=====] - 3s 7ms/step - loss: 0.0892 - a
ccuracy: 0.9731 - precision: 0.9770 - recall: 0.9693 - val_loss: 0.0492 -
val_accuracy: 0.9863 - val_precision: 0.9884 - val_recall: 0.9846
Epoch 25/100
439/439 [=====] - 3s 7ms/step - loss: 0.0863 - a
ccuracy: 0.9739 - precision: 0.9774 - recall: 0.9702 - val_loss: 0.0492 -
val_accuracy: 0.9861 - val_precision: 0.9881 - val_recall: 0.9844
Epoch 26/100
439/439 [=====] - 3s 7ms/step - loss: 0.0864 - a
ccuracy: 0.9734 - precision: 0.9771 - recall: 0.9697 - val_loss: 0.0491 -

```
val_accuracy: 0.9867 - val_precision: 0.9886 - val_recall: 0.9850
Epoch 27/100
439/439 [=====] - 3s 7ms/step - loss: 0.0884 - a
ccuracy: 0.9740 - precision: 0.9777 - recall: 0.9707 - val_loss: 0.0478 -
val_accuracy: 0.9857 - val_precision: 0.9877 - val_recall: 0.9835
Epoch 28/100
439/439 [=====] - 3s 7ms/step - loss: 0.0799 - a
ccuracy: 0.9758 - precision: 0.9791 - recall: 0.9725 - val_loss: 0.0442 -
val_accuracy: 0.9876 - val_precision: 0.9894 - val_recall: 0.9863
Epoch 29/100
439/439 [=====] - 3s 7ms/step - loss: 0.0820 - a
ccuracy: 0.9755 - precision: 0.9789 - recall: 0.9724 - val_loss: 0.0449 -
val_accuracy: 0.9867 - val_precision: 0.9891 - val_recall: 0.9856
Epoch 30/100
439/439 [=====] - 3s 7ms/step - loss: 0.0779 - a
ccuracy: 0.9748 - precision: 0.9784 - recall: 0.9718 - val_loss: 0.0441 -
val_accuracy: 0.9867 - val_precision: 0.9887 - val_recall: 0.9848
Epoch 31/100
439/439 [=====] - 3s 7ms/step - loss: 0.0730 - a
ccuracy: 0.9779 - precision: 0.9811 - recall: 0.9751 - val_loss: 0.0446 -
val_accuracy: 0.9859 - val_precision: 0.9877 - val_recall: 0.9844
Epoch 00031: early stopping
Test loss: 0.05266290158033371
Test accuracy: 0.9857142567634583
Test precision: 0.9866762161254883
Test recall: 0.9838571548461914
Test f1 score: 0.9857142857142858
Test AUC for digit: 0 0.9975965630331252
Test AUC for digit: 1 0.9964337348028355
Test AUC for digit: 2 0.9899322237602509
Test AUC for digit: 3 0.9902563737312167
Test AUC for digit: 4 0.993108146858049
Test AUC for digit: 5 0.992317975361152
Test AUC for digit: 6 0.9943650793650793
Test AUC for digit: 7 0.9950732539562421
Test AUC for digit: 8 0.9821428571428572
Test AUC for digit: 9 0.9892405775544455
```





After, we changed MaxPooling layer to AveragePooling layer. The difference between this two layers is that AveragePooling layer sums the values in the square and divides by the number of values in square. Results are worse than basic model because MaxPooling, by its characteristics, is better when we have black background, because it remembers the most white value in grey-scale.

Different number of full connected layers

```
In [44]: model_cnn_fc = keras.Sequential([
    layers.Conv2D(32, (3,3), activation="relu", input_shape=(28,28,1)),
    layers.MaxPooling2D (2,2),
    layers.Conv2D(64, (3,3), activation="relu"),
    layers.MaxPooling2D (2,2),

    layers.Flatten(),
    layers.Dense(128, activation="relu"),
    layers.Dense(32, activation="relu"),
    layers.Dropout(.5),
    layers.Dense(10, activation="softmax")
])
model_cnn_fc.summary()
predict_model(model_cnn_fc, [es], epochs=100)
```

Model: "sequential_19"

Layer (type)	Output Shape	Param #
conv2d_17 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_13 (MaxPooling)	(None, 13, 13, 32)	0
conv2d_18 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_14 (MaxPooling)	(None, 5, 5, 64)	0
flatten_19 (Flatten)	(None, 1600)	0
dense_56 (Dense)	(None, 128)	204928
dense_57 (Dense)	(None, 32)	4128
dropout_18 (Dropout)	(None, 32)	0
dense_58 (Dense)	(None, 10)	330

```
=====
Total params: 228,202
Trainable params: 228,202
Non-trainable params: 0



Epoch 1/100
439/439 [=====] - 5s 9ms/step - loss: 0.9434 - a
ccuracy: 0.6736 - precision: 0.8816 - recall: 0.5434 - val_loss: 0.0790 -
val_accuracy: 0.9759 - val_precision: 0.9819 - val_recall: 0.9717
Epoch 2/100
439/439 [=====] - 4s 8ms/step - loss: 0.2156 - a
ccuracy: 0.9378 - precision: 0.9588 - recall: 0.9172 - val_loss: 0.0600 -
val_accuracy: 0.9835 - val_precision: 0.9864 - val_recall: 0.9815
Epoch 3/100
439/439 [=====] - 3s 8ms/step - loss: 0.1334 - a
ccuracy: 0.9598 - precision: 0.9733 - recall: 0.9482 - val_loss: 0.0479 -
val_accuracy: 0.9863 - val_precision: 0.9893 - val_recall: 0.9843
Epoch 4/100
439/439 [=====] - 3s 8ms/step - loss: 0.0965 - a
ccuracy: 0.9712 - precision: 0.9809 - recall: 0.9635 - val_loss: 0.0405 -
val_accuracy: 0.9895 - val_precision: 0.9916 - val_recall: 0.9879
Epoch 5/100
439/439 [=====] - 3s 8ms/step - loss: 0.0784 - a
ccuracy: 0.9771 - precision: 0.9837 - recall: 0.9699 - val_loss: 0.0548 -
val_accuracy: 0.9866 - val_precision: 0.9877 - val_recall: 0.9853
Epoch 6/100
439/439 [=====] - 3s 8ms/step - loss: 0.0962 - a
ccuracy: 0.9722 - precision: 0.9800 - recall: 0.9644 - val_loss: 0.0438 -
val_accuracy: 0.9895 - val_precision: 0.9899 - val_recall: 0.9889
Epoch 7/100
439/439 [=====] - 3s 8ms/step - loss: 0.0635 - a
ccuracy: 0.9800 - precision: 0.9857 - recall: 0.9746 - val_loss: 0.0368 -
val_accuracy: 0.9903 - val_precision: 0.9909 - val_recall: 0.9895
Epoch 8/100
439/439 [=====] - 3s 7ms/step - loss: 0.0566 - a
ccuracy: 0.9808 - precision: 0.9863 - recall: 0.9768 - val_loss: 0.0331 -
val_accuracy: 0.9905 - val_precision: 0.9918 - val_recall: 0.9893
Epoch 9/100
439/439 [=====] - 3s 7ms/step - loss: 0.0508 - a
ccuracy: 0.9838 - precision: 0.9878 - recall: 0.9795 - val_loss: 0.0382 -
val_accuracy: 0.9916 - val_precision: 0.9922 - val_recall: 0.9909
Epoch 10/100
439/439 [=====] - 3s 8ms/step - loss: 0.0570 - a
ccuracy: 0.9811 - precision: 0.9860 - recall: 0.9768 - val_loss: 0.0343 -
val_accuracy: 0.9921 - val_precision: 0.9931 - val_recall: 0.9915
Epoch 11/100
439/439 [=====] - 3s 8ms/step - loss: 0.0483 - a
ccuracy: 0.9835 - precision: 0.9882 - recall: 0.9795 - val_loss: 0.0376 -
val_accuracy: 0.9925 - val_precision: 0.9932 - val_recall: 0.9921
Epoch 12/100
439/439 [=====] - 3s 8ms/step - loss: 0.0452 - a
ccuracy: 0.9859 - precision: 0.9890 - recall: 0.9819 - val_loss: 0.0450 -
val_accuracy: 0.9902 - val_precision: 0.9910 - val_recall: 0.9896
Epoch 13/100
439/439 [=====] - 3s 8ms/step - loss: 0.0371 - a
ccuracy: 0.9883 - precision: 0.9905 - recall: 0.9848 - val_loss: 0.0395 -
val_accuracy: 0.9918 - val_precision: 0.9922 - val_recall: 0.9909
Epoch 14/100
439/439 [=====] - 3s 8ms/step - loss: 0.0443 - a
ccuracy: 0.9852 - precision: 0.9887 - recall: 0.9817 - val_loss: 0.0439 -
val_accuracy: 0.9913 - val_precision: 0.9918 - val_recall: 0.9908
Epoch 00014: early stopping
Test loss: 0.06664026528596878
Test accuracy: 0.9888571500778198
Test precision: 0.9892780780792236
Test recall: 0.9885714054107666
Test f1 score: 0.9888571428571429
Test AUC for digit: 0 0.9990146050287297
```



```

Test AUC for digit: 1 0.9951085363883655
Test AUC for digit: 2 0.9946951426069947
Test AUC for digit: 3 0.9941206967376037
Test AUC for digit: 4 0.9936140097241216
Test AUC for digit: 5 0.9941211483938233
Test AUC for digit: 6 0.9940476190476191
Test AUC for digit: 7 0.9912349278979704
Test AUC for digit: 8 0.9904761904761904
Test AUC for digit: 9 0.9916870128535711

```



The performance of a published network (LeNet5, VGG, Yolo, etc) for recognizing MNIST Digits

We decided to implement the architecture of LeNet5 network. The LeNet-5 architecture consists of two sets of convolutional and average pooling layers, followed by a flattening convolutional layer, then two fully-connected layers and finally a softmax classifier.

LeNet5

This is an implementation of LeNet5 (slightly different, because input shape is 28x28 and in original version was 32x32). Despite of its age the model is pretty accurate (with accuracy 98.9%). This is close to our best models, and it does not have too big number of parameters (only 60,074).

```
In [66]: lenet5 = keras.Sequential([
```

```

layers.Conv2D(filters=6, kernel_size=(3, 3), activation='relu', input_shape=(28, 28, 1)),
layers.AveragePooling2D(),

layers.Conv2D(filters=16, kernel_size=(3, 3), activation='relu'),
layers.AveragePooling2D(),

layers.Flatten(),

layers.Dense(units=120, activation='relu'),
layers.Dense(units=84, activation='relu'),
layers.Dense(units=10, activation='softmax')
])

lenet5.summary()
predict_model(lenet5, [es], epochs=100)

```

Model: "sequential_27"

Layer (type)	Output Shape	Param #
conv2d_30 (Conv2D)	(None, 26, 26, 6)	60
average_pooling2d_12 (AveragePooling2D)	(None, 13, 13, 6)	0
conv2d_31 (Conv2D)	(None, 11, 11, 16)	880
average_pooling2d_13 (AveragePooling2D)	(None, 5, 5, 16)	0
flatten_27 (Flatten)	(None, 400)	0
dense_82 (Dense)	(None, 120)	48120
dense_83 (Dense)	(None, 84)	10164
dense_84 (Dense)	(None, 10)	850

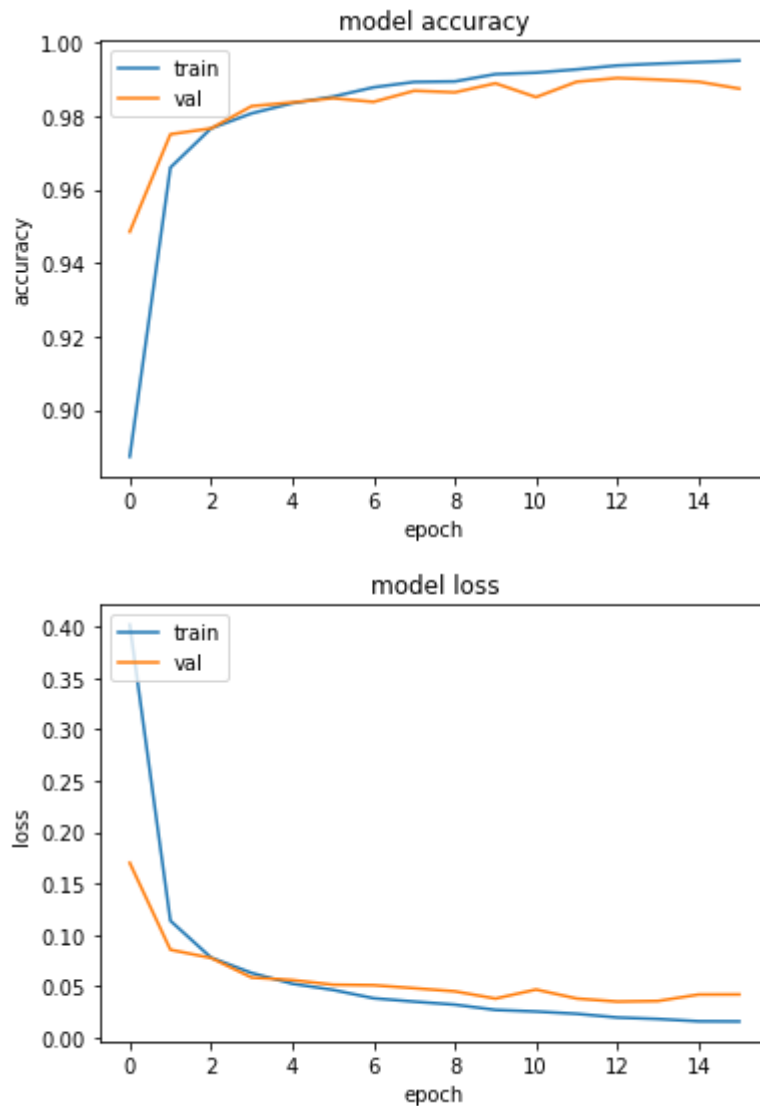
Total params: 60,074
 Trainable params: 60,074
 Non-trainable params: 0

```

Epoch 1/100
439/439 [=====] - 4s 7ms/step - loss: 0.8082 - accuracy: 0.7743 - precision: 0.8916 - recall: 0.6138 - val_loss: 0.1701 - val_accuracy: 0.9485 - val_precision: 0.9569 - val_recall: 0.9416
Epoch 2/100
439/439 [=====] - 3s 7ms/step - loss: 0.1352 - accuracy: 0.9603 - precision: 0.9669 - recall: 0.9535 - val_loss: 0.0855 - val_accuracy: 0.9750 - val_precision: 0.9798 - val_recall: 0.9723
Epoch 3/100
439/439 [=====] - 3s 7ms/step - loss: 0.0792 - accuracy: 0.9762 - precision: 0.9797 - recall: 0.9734 - val_loss: 0.0775 - val_accuracy: 0.9766 - val_precision: 0.9808 - val_recall: 0.9743
Epoch 4/100
439/439 [=====] - 3s 7ms/step - loss: 0.0660 - accuracy: 0.9803 - precision: 0.9830 - recall: 0.9780 - val_loss: 0.0585 - val_accuracy: 0.9827 - val_precision: 0.9855 - val_recall: 0.9812
Epoch 5/100
439/439 [=====] - 3s 7ms/step - loss: 0.0530 - accuracy: 0.9835 - precision: 0.9857 - recall: 0.9815 - val_loss: 0.0560 - val_accuracy: 0.9837 - val_precision: 0.9865 - val_recall: 0.9815
Epoch 6/100
439/439 [=====] - 3s 7ms/step - loss: 0.0451 - accuracy: 0.9858 - precision: 0.9874 - recall: 0.9846 - val_loss: 0.0515 - val_accuracy: 0.9848 - val_precision: 0.9877 - val_recall: 0.9838
Epoch 7/100
439/439 [=====] - 3s 7ms/step - loss: 0.0382 - accuracy: 0.9878 - precision: 0.9892 - recall: 0.9867 - val_loss: 0.0510 -

```

val_accuracy: 0.9838 - val_precision: 0.9872 - val_recall: 0.9825
Epoch 8/100
439/439 [=====] - 3s 7ms/step - loss: 0.0330 - accuracy: 0.9898 - precision: 0.9910 - recall: 0.9890 - val_loss: 0.0480 - val_accuracy: 0.9869 - val_precision: 0.9881 - val_recall: 0.9848
Epoch 9/100
439/439 [=====] - 3s 7ms/step - loss: 0.0357 - accuracy: 0.9880 - precision: 0.9894 - recall: 0.9870 - val_loss: 0.0450 - val_accuracy: 0.9864 - val_precision: 0.9880 - val_recall: 0.9848
Epoch 10/100
439/439 [=====] - 3s 7ms/step - loss: 0.0265 - accuracy: 0.9913 - precision: 0.9920 - recall: 0.9902 - val_loss: 0.0379 - val_accuracy: 0.9889 - val_precision: 0.9899 - val_recall: 0.9879
Epoch 11/100
439/439 [=====] - 3s 7ms/step - loss: 0.0245 - accuracy: 0.9916 - precision: 0.9922 - recall: 0.9911 - val_loss: 0.0468 - val_accuracy: 0.9851 - val_precision: 0.9865 - val_recall: 0.9843
Epoch 12/100
439/439 [=====] - 3s 7ms/step - loss: 0.0253 - accuracy: 0.9920 - precision: 0.9927 - recall: 0.9915 - val_loss: 0.0380 - val_accuracy: 0.9893 - val_precision: 0.9905 - val_recall: 0.9883
Epoch 13/100
439/439 [=====] - 3s 7ms/step - loss: 0.0206 - accuracy: 0.9935 - precision: 0.9939 - recall: 0.9928 - val_loss: 0.0350 - val_accuracy: 0.9903 - val_precision: 0.9913 - val_recall: 0.9887
Epoch 14/100
439/439 [=====] - 3s 7ms/step - loss: 0.0168 - accuracy: 0.9947 - precision: 0.9953 - recall: 0.9943 - val_loss: 0.0356 - val_accuracy: 0.9899 - val_precision: 0.9909 - val_recall: 0.9893
Epoch 15/100
439/439 [=====] - 3s 7ms/step - loss: 0.0133 - accuracy: 0.9955 - precision: 0.9959 - recall: 0.9952 - val_loss: 0.0419 - val_accuracy: 0.9893 - val_precision: 0.9900 - val_recall: 0.9883
Epoch 16/100
439/439 [=====] - 3s 7ms/step - loss: 0.0129 - accuracy: 0.9959 - precision: 0.9962 - recall: 0.9957 - val_loss: 0.0420 - val_accuracy: 0.9874 - val_precision: 0.9884 - val_recall: 0.9870
Restoring model weights from the end of the best epoch.
Epoch 00016: early stopping
Test loss: 0.039734747260808945
Test accuracy: 0.9887142777442932
Test precision: 0.9896981120109558
Test recall: 0.9881428480148315
Test f1 score: 0.9887142857142858
Test AUC for digit: 0 0.9967844920645857
Test AUC for digit: 1 0.9953059136050564
Test AUC for digit: 2 0.9949662450022941
Test AUC for digit: 3 0.9916327889937758
Test AUC for digit: 4 0.9944156815297652
Test AUC for digit: 5 0.9953975922796036
Test AUC for digit: 6 0.9949879467017794
Test AUC for digit: 7 0.9939484045292735
Test AUC for digit: 8 0.986462230229477
Test AUC for digit: 9 0.9932852101356561



The best network

```
In [117]: from sklearn.metrics import confusion_matrix

y_pred = model_cnn_long.predict(x_test)
y_pred1 = list(np.argmax(y_pred, axis=1))
y_test1 = list(np.argmax(y_test, axis = 1))

confusion_matrix = confusion_matrix(y_test1, y_pred1)

print(confusion_matrix)

[[654  0  0  0  0  0  0  0  0  0]
 [  0 765  0  0  0  0  0  1  0  1]
 [  1  1 697  1  0  0  0  3  1  0]
 [  0  0  0 707  0  0  0  2  0  0]
 [  0  0  0  0 670  0  1  0  0  2]
 [  1  1  0  0  0 648  0  0  0  2]
 [  1  0  0  1  0  1 697  0  0  0]
 [  0  0  0  0  0  1  0 745  0  0]
 [  0  0  0  1  0  0  0  0 683  2]
 [  0  1  0  0  3  1  0  1  0 703]]
```

We chose model with best accuracy. It was a model with 4 convolutional layers and this is confusion matrix of this model.

In []:

